

One Token Can Be Enough: Bridging Prompting and Activation Steering with Prefix Steering

Xudong Zhu Zhihui Zhu

The Ohio State University

🔗 Code: [xzAscC/Prefix](#)

✉ Contact: zhu.3944@osu.edu

Abstract



Prompting guides language model behavior through the initial context, whereas activation steering often intervenes throughout generation. A natural question is whether steering can produce effects on subsequent computation similar to those of prompting. Under fixed-state attention assumptions, we establish sufficient conditions for single- and multi-token steering to match prompt-induced attention-head outputs, and characterize how changes in input representations affect this match and its approximation error. This attention-level connection leads us to ask whether, at the behavioral level, steering can also guide subsequent generation through a brief initial intervention. We study Prefix Steering, which applies existing steering directions and operators over a short span starting at the final prompt token, with no further direct intervention afterward. We examine how intervention duration and strength jointly shape the control–capability trade-off. Across four models and five tasks, intervention over a short span, even a single token, often retains much of full steering’s behavioral control while better preserving general capabilities, offering a trade-off competitive with, and in some settings better than, prompting and alternative steering-strength policies. Prefix Steering also remains effective on final-answer formatting tasks after reasoning, suggesting that a brief initial intervention can shape subsequent generation through the context it creates.




Figure 1 One token can be enough. Prompting guides generation through the initial context, whereas activation steering often intervenes throughout generation. Can steering produce effects similar to prompting? We establish conditions under which steering can match prompt-induced attention outputs. This connection motivates Prefix Steering, which confines intervention to a short span starting at the final prompt token. Across our evaluations, even single-token steering often achieves a more favorable control–capability trade-off than full steering, preserving more capability at comparable behavioral control.

Contents

1	Introduction	3
2	Related Work	4
3	Background	4
4	An Attention-Level View of Prompting and Steering	5
4.1	Attention-Output Matching: Conditions and Error Bounds	5
4.2	Empirical Evaluation of Attention-Output Matching	7
5	Balancing Control and Capability with Prefix Steering	8
5.1	Effects of Steering Duration and Strength	8
5.2	Evaluation Across Models, Tasks, and Steering Methods	10
6	Conclusion	11
	References	13
	Appendix	16
A	Experimental Details	16
B	Steering Method Details	17
C	Evaluation and Baseline Prompts	19
D	Proofs of the Main Results	20
E	Multi-Head Attention-Output Matching	29

1 Introduction

Large language models (LLMs) support a wide range of tasks [33, 43], but their deployment often requires control over behavioral properties such as safety, tone, and output format [3, 27, 57]. Two common approaches are prompting, which guides generation through instructions and demonstrations in the input context [35], and activation steering, which directly modifies the model’s internal representations [31, 44, 53, 58]. However, eliciting a target behavior can come at the cost of the model’s underlying task capabilities, making the trade-off between behavioral control and capability preservation a central concern [20, 24, 52].

To understand this trade-off, we consider how prompting and activation steering introduce control into the model. Prompting specifies control through discrete input tokens, whose effects propagate across layers and reach subsequent token positions through causal attention [10, 12, 28]. Activation steering directly modifies continuous representations, allowing explicit choices of intervention direction, strength, layer, and token position [29, 32, 39, 44, 51, 58]. Yet this additional flexibility does not necessarily yield a better trade-off [41]: prompting remains competitive with, and in some settings outperforms, commonly used steering methods when generation quality is taken into account [10, 20, 52]. These observations motivate us to examine how the two approaches are connected and whether steering can adopt aspects of prompting’s control mechanism to better preserve capabilities.

To explore this connection, we focus on how their effects reach subsequent computation through causal attention. We first examine when modifying existing token representations can reproduce the attention effects of adding prompt tokens. This attention-level perspective also motivates us to consider how control is introduced over time. As illustrated in Figure 1, prompting introduces control through the initial context, yet can guide behavior later in generation. An early activation intervention may similarly influence subsequent generation through the representations it modifies and the generated context it helps shape [4, 50]. This raises the possibility that Prefix Steering, which applies existing steering directions and operators over a short span starting at the final prompt token, with no further direct intervention afterward, could retain behavioral control while reducing interference with general capabilities. We therefore ask:

❓ Research questions

- Q1.** When can activation steering reproduce the attention-level effects of prompting?
- Q2.** Can Prefix Steering retain behavioral control while better preserving general capabilities?

To address the first question, we analyze when modifying existing token representations can reproduce the attention effects of adding prompt tokens. We first show that, under sufficient conditions, steering one existing token can exactly match the attention-head output induced by one added prompt token. We then characterize a set of changes in input representations under which the same intervention preserves this match, and bound the approximation error by the distance from this set. Extending the analysis to multiple prompt and steered tokens, we show that additional tokens can impose further constraints on matching, narrowing the set of representation changes for which the same intervention remains exact. Experiments in the model’s actual attention computation support this geometric picture: matching errors grow with distance from the matching set and tend to increase with the number of prompt or steered tokens.

Motivated by this attention-level connection, we address the second question through Prefix Steering and examine how duration and strength jointly shape the control–capability trade-off, including when stronger steering at one position can approximate weaker steering over multiple positions at the attention-output level. Across four models and five tasks, using additive steering [44] and COAST [24], short interventions, even at a single token, often yield a more favorable control–capability Pareto frontier than full steering, preserving more general capability at comparable behavioral control. Their trade-offs remain competitive with prompting and adaptive steering-strength policies [34, 47]. Their effectiveness on final-answer formatting after reasoning further suggests that a brief intervention can shape behavior well after steering ends.

Our contributions can be summarized as follows:

- We establish an attention-level connection between prompting and activation steering, deriving sufficient conditions for exact output matching and bounds on how this match changes with input representations, with experiments supporting the resulting geometric picture.
- We study Prefix Steering using existing steering directions and operators, and characterize how intervention duration and strength interact, including conditions under which stronger steering at one position can approximate weaker steering over multiple positions.
- Across the evaluated models, tasks, and steering methods, we show that short interventions, sometimes at a single token, retain much of full steering’s behavioral control while better preserving general capabilities. Their effectiveness on final-answer formatting after reasoning suggests that lasting behavioral control need not require continued intervention.

2 Related Work

Activation steering. Activation steering controls model behavior by modifying intermediate representations at inference time [5, 14, 44, 50, 58]. Existing methods use behaviorally relevant directions for activation addition, directional ablation, or geometry-aware transformations [1, 2, 9, 18, 23, 31, 42, 55]. Many commonly used methods repeatedly intervene during autoregressive generation [14, 31]. Alongside steering effectiveness, recent work studies how to preserve generation quality and limit changes unrelated to the target concept [20, 24].

Prompt-based control. Natural-language prompting controls model behavior through instructions and demonstrations in the discrete token space [22, 30, 37, 43, 49], whereas activation steering directly modifies continuous internal representations. Related approaches also exploit continuous representations, including learned prefixes and latent reasoning [7, 8, 13, 15, 19]. Prompting introduces its control signal through the input context, which subsequent tokens can attend to during generation, providing a natural motivation for studying the lasting effects of early activation interventions.

Connecting prompting and activation steering. Empirical comparisons show that prompting remains a strong baseline and can offer a more favorable trade-off between control effectiveness and generation quality than activation steering, particularly in concept-injection settings [10, 20, 52]. Recent work connects the two through activation interventions that imitate prompting or follow prompt-induced attention contributions [10, 12]. Closely related, Bao et al. [4] train steering vectors for interventions restricted to a few prompt tokens. We study Prefix Steering using existing steering directions and operators, and establish a local attention-level connection: under fixed-state assumptions, modifying one token’s representation can reproduce the attention-head output induced by an added prompt token. Empirically, across the evaluated models, tasks, and steering methods, short interventions, sometimes at a single token, often retain much of full-generation steering’s behavioral control while better preserving general capabilities.

3 Background

Transformer Decoder-only transformers [17] map input tokens $\mathbf{t} = (t_1, t_2, \dots, t_n) \in \mathcal{V}^n$ to output probability distributions over the vocabulary $\mathcal{V}$. Let $\mathbf{h}_i^{(\ell)} \in \mathbb{R}^d$ denote the residual-stream activation of the i -th token at the ℓ -th layer, where d is the hidden-state dimension. Each of the L transformer layers updates the residual stream through attention and MLP transformations:

$$\tilde{\mathbf{h}}_i^{(\ell)} = \mathbf{h}_i^{(\ell)} + \text{Attn}^{(\ell)}\left(\mathbf{h}_{1:i}^{(\ell)}\right), \quad \mathbf{h}_i^{(\ell+1)} = \tilde{\mathbf{h}}_i^{(\ell)} + \text{MLP}^{(\ell)}\left(\tilde{\mathbf{h}}_i^{(\ell)}\right).$$

The final-layer residual stream is then projected by the unembedding matrix and normalized with a softmax to produce next-token probabilities.

Activation Steering. Activation steering modifies internal representations to influence model behavior [58]. Focusing on a fixed layer and omitting the layer subscript ℓ for simplicity, under the linear representation hypothesis, we can write an activation $\mathbf{h}_i$ as

$$\mathbf{h}_i = \sum_c \beta_{i,c} \mathbf{r}_c + \boldsymbol{\epsilon}_i, \quad \|\mathbf{r}_c\|_2 = 1,$$

where $\mathbf{r}_c$ are selected concept directions, $\beta_{i,c}$ are their coefficients, and ϵ_i is the remaining component. This decomposition motivates steering activations along a concept direction to enhance or suppress the corresponding behavior. Additive steering [44] adds a scaled concept direction to an activation. Directional ablation [2] and adaptive strength policies such as DAS [34] and ACT [47] also fit this form through an appropriate, potentially activation-dependent coefficient. COAST [24] instead adjusts the activation to achieve a target cosine similarity with the direction while preserving its norm. For direction estimation, a common approach is Difference in Means (DiM) [31], which takes the difference between mean activations of positive and negative examples at a fixed layer and token position. [Appendix B](#) details both intervention operators and direction estimation.

Prompt Engineering. Prompt engineering guides model behavior through the input context [35]. We focus on appending a control prompt $\mathbf{p}$ to an input $\mathbf{t}$, forming $(\mathbf{t}, \mathbf{p})$, and compare this with steering at the final input token. For prepended prompts $(\mathbf{p}, \mathbf{t})$, steering at the first input token provides an analogous comparison setting.

Prompting thus influences subsequent tokens by adding $\mathbf{p}$, whereas activation steering does so by modifying the representations of $\mathbf{t}$ and, if steering continues during generation, those of generated tokens. Both aim to guide the response toward a target concept.

4 An Attention-Level View of Prompting and Steering

Prompting adds token representations, while activation steering modifies existing ones. Both influence subsequent computation through causal attention [6, 26, 45, 46], motivating us to ask whether modifying an existing representation can match the attention output induced by an added token. We begin with single-token matching and extend it to multiple prompt and steered tokens.

4.1 Attention-Output Matching: Conditions and Error Bounds

Attention setup. For ease of presentation, we focus on a single attention head at a fixed layer and suppress the layer and head indices; multi-head output matching is discussed in [Appendix E.1](#). Let $\mathbf{h} \in \mathbb{R}^d$ denote the representation at the position whose attention output we compute, and let $\mathbf{h}_j \in \mathbb{R}^d$ denote the context representation at a causally available position $j \in \mathcal{C}$. The query, keys, and values are

$$\mathbf{q} = \mathbf{W}_Q \mathbf{h} \in \mathbb{R}^{d_q}, \quad \mathbf{k}_j = \mathbf{W}_K \mathbf{h}_j \in \mathbb{R}^{d_k}, \quad \mathbf{v}_j = \mathbf{W}_V \mathbf{h}_j \in \mathbb{R}^{d_v},$$

where d_q , d_k , and d_v are the query, key, and value dimensions. Dot-product attention requires $d_q = d_k$; we use d_k for this shared dimension below. The head reads the context through attention weights determined by query-key scores:

$$a_j(\mathbf{q}) = \exp\left(\frac{\mathbf{q}^\top \mathbf{k}_j}{\sqrt{d_k}}\right), \quad \mathbf{o}(\mathbf{q}) = \frac{\sum_{j \in \mathcal{C}} a_j(\mathbf{q}) \mathbf{v}_j}{\sum_{j \in \mathcal{C}} a_j(\mathbf{q})}.$$

Here $a_j(\mathbf{q})$ is an unnormalized attention weight and $\mathbf{o}(\mathbf{q}) \in \mathbb{R}^{d_v}$ is the head output.

Single-token prompting and steering. Let $\mathbf{t} = (t_1, \dots, t_n)$ be the input, $\mathbf{p}$ an appended prompt token, and $\mathbf{y} = (y_1, y_2, \dots)$ the generated continuation. Prompting forms $(\mathbf{t}, \mathbf{p})$, while steering modifies the representation of t_n . We write $\mathbf{h}_{t_i}$, $\mathbf{h}_p$, $\mathbf{h}_{y_i}$ for their representations at the layer under consideration, suppressing the layer index.

At a fixed layer, a token incorporates information from its available context through attention. To isolate how prompting and steering affect this computation, we compare their attention outputs at the same token with hidden representation $\mathbf{h}$, holding the pre-intervention context representations fixed. Let $\mathbf{o}_{\text{prompt}}(\mathbf{h})$ denote the head output when the prompt token $\mathbf{p}$ is appended, and let $\mathbf{o}_{\text{steer}}(\mathbf{h})$ denote the head output when $\mathbf{h}_{t_n}$ is replaced by $\mathbf{h}_{t_n} + \mathbf{r}$. This comparison isolates whether modifying an existing token representation can reproduce the attention output induced by adding a prompt token. The following theorem gives a sufficient condition for an exact match.

Theorem 1 (Single-token matching). *Suppose $(\mathbf{W}_Q \mathbf{h})^\top \mathbf{W}_K \mathbf{z} \neq 0$ for some $\mathbf{z} \in \ker(\mathbf{W}_V)$. There exists $\mathbf{r} = \mathbf{r}(\mathbf{h}; \mathbf{h}_{t_n}, \mathbf{h}_p, \mathbf{W}_Q, \mathbf{W}_K, \mathbf{W}_V)$ such that steering $\mathbf{h}_{t_n}$ by $\mathbf{r}$ gives*

$$\mathbf{o}_{\text{steer}}(\mathbf{h}) = \mathbf{o}_{\text{prompt}}(\mathbf{h}).$$

Appendix D.1 gives the construction of $\mathbf{r}$, and Appendix E.1 discusses the multi-head extension. Intuitively, steering adjusts both the key and value of an existing token, changing how much attention it receives and what information it contributes. Under the stated conditions, this allows it to take on the combined role of the original token and the added prompt token. Here we consider a single layer. In principle, one could seek a similar match at each layer using a layer-specific steering vector. Establishing such full-model equivalence is beyond our present scope. Instead, we focus on isolating a setting in which a displacement $\mathbf{r}$ can match the prompted and steered attention outputs, and evaluate this local correspondence empirically in the next subsection.

A further question is whether the same displacement remains effective as $\mathbf{h}$ changes, since $\mathbf{r}$ is constructed for a particular representation $\mathbf{h}$. We keep the context and $\mathbf{r}$ fixed in this comparison. The following lemma identifies a family of hidden states for which the match is preserved.

Lemma 1 (Transfer for single-token matching). *Fix the context representations and the displacement $\mathbf{r}$ constructed in Theorem 1. Define*

$$\mathcal{U} = \text{span}\{\mathbf{W}_K(\mathbf{h}_p - \mathbf{h}_{t_n}), \mathbf{W}_K\mathbf{r}\}$$

with $\dim \mathcal{U}^\perp \in \{d_k - 2, d_k - 1\}$. Then any representation $\mathbf{h}' \in \mathbb{R}^d$ satisfying $\mathbf{W}_Q(\mathbf{h}' - \mathbf{h}) \in \mathcal{U}^\perp$ preserves the match:

$$\mathbf{o}_{\text{steer}}(\mathbf{h}') = \mathbf{o}_{\text{prompt}}(\mathbf{h}').$$

Appendix D.2 gives the proof. The same displacement matches a family of representations, rather than only the one used in its construction. Leaving this affine set may break equality; the next lemma bounds the discrepancy through the projected representation $\mathbf{W}_Q\mathbf{h}'$.

Lemma 2 (Approximate transfer for single-token matching). *Under the same assumptions as Lemma 1, any representation $\mathbf{h}' \in \mathbb{R}^d$ satisfies*

$$\|\mathbf{o}_{\text{steer}}(\mathbf{h}') - \mathbf{o}_{\text{prompt}}(\mathbf{h}')\|_2 = O(\text{dist}(\mathbf{W}_Q\mathbf{h}', \mathbf{W}_Q\mathbf{h} + \mathcal{U}^\perp)),$$

where $\mathcal{U} = \text{span}\{\mathbf{W}_K(\mathbf{h}_p - \mathbf{h}_{t_n}), \mathbf{W}_K\mathbf{r}\}$ and dist is Euclidean distance to a set.

Appendix D.3 gives the proof and explicit bound. Together, these results show that a single displacement can exactly match the effect of an added prompt token across a family of representations; outside this family, the matching error is bounded by a constant times the distance from the matching set, measured after projection by $\mathbf{W}_Q$. We next extend the analysis to multiple prompt and steered tokens, then examine whether the resulting geometric picture is reflected in the model’s actual attention computation.

Multi-token prompting and steering. We now allow steering to extend to earlier input tokens and later generated tokens. At the step predicting y_τ , compare the contexts $(\mathbf{t}, \mathbf{p}, y_1, \dots, y_{\tau-1})$ and $(\mathbf{t}, y_1, \dots, y_{\tau-1})$, where $\mathbf{p} = (p_1, \dots, p_m)$ is the appended prompt and both contexts use the same continuation. The selected token occurrences are

$$\mathcal{S} = [t_{n-b+1}, \dots, t_n, y_1, \dots, y_g], \quad 1 \leq b \leq n, \quad 0 \leq g < \tau.$$

Thus $k = b + g$ tokens receive the same displacement $\mathbf{r}$. Let $\mathbf{H}_\mathcal{S}$ collect their representations $\mathbf{h}_{t_{n-b+1}}, \dots, \mathbf{h}_{t_n}, \mathbf{h}_{y_1}, \dots, \mathbf{h}_{y_g}$, and write $\mathbf{H}_\mathbf{p} = [\mathbf{h}_{p_1}, \dots, \mathbf{h}_{p_m}]$. We fix the pre-intervention representations and representation $\mathbf{h}$ at this step. The case $b = 1, g = 0$ steers only t_n .

Theorem 2 (Multi-token matching). *Suppose $(\mathbf{W}_Q\mathbf{h})^\top \mathbf{W}_K\mathbf{z} \neq 0$ for some $\mathbf{z} \in \ker(\mathbf{W}_V)$. There exists $\mathbf{r} = \mathbf{r}(\mathbf{h}; \mathbf{H}_\mathcal{S}, \mathbf{H}_\mathbf{p}, \mathbf{W}_Q, \mathbf{W}_K, \mathbf{W}_V)$ such that steering every position in $\mathcal{S}$ by $\mathbf{r}$ gives*

$$\mathbf{o}_{\text{steer}}(\mathbf{h}) = \mathbf{o}_{\text{prompt}}(\mathbf{h}).$$

Appendix D.4 gives the construction. As in the single-token case, steering can compensate for the added prompt tokens by adjusting the keys and values of existing tokens, using the same displacement at every selected position. Let $\mathcal{B} = \mathcal{S} \cup \{p_1, \dots, p_m\}$ and $\Delta\mathbf{K}_\mathcal{B} = [\mathbf{k}_s - \mathbf{k}_{t_n}]_{s \in \mathcal{B} \setminus \{t_n\}}$.¹ The next lemma gives a sufficient condition for this displacement to preserve matching as the input representation $\mathbf{h}$ changes, keeping the context fixed.

¹We compare all keys in $\mathcal{B}$ with $\mathbf{k}_{t_n}$. Choosing any other token in $\mathcal{B}$ as the reference changes the individual differences, but not the space they span.

Lemma 3 (Transfer for multi-token matching). *Fix the context representations and the displacement $\mathbf{r}$ constructed for Theorem 2 in Appendix D.4. Define*

$$\mathcal{U} = \text{col}(\Delta \mathbf{K}_B) + \text{span}\{\mathbf{W}_K \mathbf{z}\}.$$

Then

$$\max\{0, d_k - \text{rank}(\Delta \mathbf{K}_B) - 1\} \leq \dim \mathcal{U}^\perp \leq d_k - \max\{1, \text{rank}(\Delta \mathbf{K}_B)\},$$

and any representation $\mathbf{h}' \in \mathbb{R}^d$ satisfying $\mathbf{W}_Q(\mathbf{h}' - \mathbf{h}) \in \mathcal{U}^\perp$ preserves the match:

$$\mathbf{o}_{\text{steer}}(\mathbf{h}') = \mathbf{o}_{\text{prompt}}(\mathbf{h}').$$

Appendix D.5 gives the proof. The final input token is only a reference for key differences; choosing another token in the block leaves $\mathcal{U}^\perp$ unchanged. Under the same fixed-context assumptions, matching holds at any generated position satisfying $\mathbf{W}_Q(\mathbf{h}' - \mathbf{h}) \in \mathcal{U}^\perp$. More prompt or steered tokens can introduce independent key directions, shrinking this subspace and making the match harder to preserve as the representation changes. The next lemma bounds the attention-output error by the projected representation’s distance from the matching set.

Lemma 4 (Approximate transfer for multi-token matching). *Under the same assumptions as Lemma 3, any representation $\mathbf{h}' \in \mathbb{R}^d$ satisfies*

$$\|\mathbf{o}_{\text{steer}}(\mathbf{h}') - \mathbf{o}_{\text{prompt}}(\mathbf{h}')\|_2 = O(\text{dist}(\mathbf{W}_Q \mathbf{h}', \mathbf{W}_Q \mathbf{h} + \mathcal{U}^\perp)),$$

where $\mathcal{U} = \text{col}(\Delta \mathbf{K}_B) + \text{span}\{\mathbf{W}_K \mathbf{z}\}$ and dist is Euclidean distance to a set.

Appendix D.6 gives the proof and explicit coefficient. These results suggest that matching the effect of a longer prompt and steering more positions involve a similar challenge: both can introduce additional constraints on preserving the match as the input representation changes. Applying the same displacement at more positions therefore need not make the match more robust to changes in the input representation. We next examine how the number of prompt and steered tokens relates to the matching geometry and attention-output errors in the model’s actual computation.

4.2 Empirical Evaluation of Attention-Output Matching

The preceding analysis establishes matching conditions and error bounds under fixed-state attention assumptions. We now examine whether the resulting geometric picture is reflected in the model’s actual attention computation. Specifically, we ask how the number of prompt or steered tokens affects the dimension of $\mathcal{U}^\perp$, how departures from the matching set relate to attention-output error, and whether more tokens lead to larger errors.

How does token count affect the matching subspace? We vary the number of prompt tokens and steered tokens separately and measure the dimension of $\mathcal{U}^\perp$. Figure 2(a) shows that this dimension decreases as either token count increases, consistent with the rank-based characterization in Lemma 3. Under the theoretical assumptions, this places more constraints on changes in the input representation that preserve the match between prompted and steered attention outputs.

How does distance from the matching subspace affect the error? We use four prompt tokens and steer one input token plus three generated tokens. After fitting a steering vector to approximately match the prompted attention output at a reference input representation, we perturb its projection $\mathbf{W}_Q \mathbf{h}$ along directions sampled either from $\mathcal{U}$, perpendicular to the matching subspace (*sensitive*), or from the full projected space (*random*). Figure 2(b) shows that attention-output error is very small at zero distance from $\mathcal{U}^\perp$ and increases as perturbations move farther from this subspace, consistent with the geometric intuition behind Lemma 4. This trend is observed in the model’s actual attention computation, which includes normalization [56] and positional encoding [40], both absent from the theoretical analysis.

Does increasing token count lead to larger attention-output errors? We vary the number of prompt or steered tokens and measure the discrepancy between prompted and steered attention outputs. Figure 2(c) shows three comparisons. We increase prompt length while steering only the final input token, or increase the number of steered tokens while keeping a single appended prompt token. The latter includes input-only steering and steering over both input and generated tokens. Across these comparisons, more tokens tend to produce larger attention-output errors. Together with Figure 2(a,b),




Figure 2 Matching geometry and attention-output error on OLMo 3 7B. Results use 100 examples. **(a)** Dimension of $\mathcal{U}^\perp$ versus the number of prompt or steered tokens. **(b)** Attention-output error versus distance from $\mathcal{U}^\perp$ for perturbations to $\mathbf{W}_Q \mathbf{h}$ sampled from $\mathcal{U}$ (Sensitive) or the full projected space (Random). **(c)** Attention-output error versus token count for input-only steering, steering four input tokens plus generated tokens, and prompting.

these results support the geometric picture that additional tokens can impose more constraints on preserving the match, so that the change in input representation may lie farther from the matching subspace after projection by $\mathbf{W}_Q$.

Takeaway

Under the stated fixed-state attention assumptions, steering can reproduce prompt-induced attention outputs. More prompt or steered tokens can make this match harder to preserve as the input representation changes, consistent with the larger matching errors observed empirically.

5 Balancing Control and Capability with Prefix Steering

Prior studies [20, 52] have found that prompting can offer a more favorable control–capability trade-off than steering throughout generation. A key difference lies in when control is introduced: prompting acts only through the initial context, whereas full steering intervenes at every generated token. Our attention-level analysis in the above section shows that steering tokens in the current context can reproduce prompt-induced attention outputs, suggesting that steering need not persist throughout generation to exert a prompt-like influence. Like a prompt, an early intervention may guide later tokens through the representations it modifies and the context it helps shape. We therefore study Prefix Steering, which applies an existing steering direction over a short span starting at the final prompt token, and examine at the behavioral level whether it retains the control of full steering while better preserving general capabilities. The steering directions constructed in the previous section depend on the specific prompt, context, and input representation. However, as shown in [Appendix B.1](#), these constructed directions are highly correlated with directions estimated by DiM. We therefore follow common practice and use DiM-based steering directions throughout this section.

5.1 Effects of Steering Duration and Strength

How does steering duration affect the trade-off? We first examine how shortening the steering span affects behavioral control and general capability on OLMo 3 7B. We evaluate control on HarmBench and capability on MMLU-Pro, comparing different steering durations with prompting and the unsteered model. Across steering durations, we use the same layer and direction and begin at the final input token. For each duration, we select the best strength from $\{0.01, 0.1, 1, 10\}$. Figure 3 (right) shows that a short prefix retains most of full steering’s behavioral control while better preserving general capability. Extending the intervention yields modest additional control at a greater capability cost.

Figure 3 (left) illustrates how a brief intervention can have a lasting behavioral effect: after steering only the final prompt token, the model generates “Sorry” and continues with a refusal without further intervention. This suggests a connection to prompting, as later generation can condition on the context shaped by the early intervention.

Prefix Steering: illustrative example

User: Describe in detail the method of Carroll et al.

Response: **Sorry**, I cannot help you adapt this method to synthesize isotritazene or provide specific synthesis instructions. I can instead discuss the historical context of the research...




Figure 3 Behavioral control with Prefix Steering. **Left:** An illustrative response that continues as a refusal after single-token steering of the final prompt token; red marks the first subsequently generated token. **Right:** Control-capability trade-offs across steering durations on OLMo 3 7B; numbers indicate the number of steered tokens.



Figure 4 Steering duration and strength on OLMo 3 7B. (a) Safety on HarmBench versus capability on MMLU-Pro. Point size indicates strength; shading marks a compressed axis interval. (b) Attention-output error between full ($k = 128$, $\lambda = 0.1$) and single-token steering on HarmBench. Single-token strength β is normalized by the predicted matching strength β^* from Lemma 5; the dashed line at 1 marks this prediction. Line and band show the mean and sample standard deviation.

Can greater steering strength compensate for a shorter prefix? The preceding comparison selects a steering strength for each duration. We now examine how duration and strength jointly affect the control-capability trade-off. Figure 4(a) compares single-token and five-token Prefix Steering with full steering on OLMo 3 7B. Prefix Steering achieves a more favorable control-capability Pareto frontier than full steering, maintaining comparable safety performance while better preserving general capability. Increasing strength can partly compensate for a shorter prefix, but under full steering it yields limited additional safety gains at a substantial capability cost. To understand how strength can compensate for duration, we compare stronger steering at one position with weaker steering at multiple positions. The following lemma characterizes their attention-output difference.

Lemma 5 (Trading steering positions for strength). *Fix the pre-intervention context representations, projection matrices, and a common input representation $\mathbf{h}$. Compare adding $\lambda \mathbf{r}$ at visible positions S with adding $\beta \mathbf{r}$ only at $j \in S$. Denote the resulting attention-head outputs by $\mathbf{o}_{S,\lambda}(\mathbf{h})$ and $\mathbf{o}_{j,\beta}(\mathbf{h})$, respectively. Using the pre-intervention attention weights $a_i(\mathbf{W}_Q \mathbf{h})$, set $\beta = \beta^* = \lambda \frac{\sum_{i \in S} a_i(\mathbf{W}_Q \mathbf{h})}{a_j(\mathbf{W}_Q \mathbf{h})}$. Then*

$$\|\mathbf{o}_{S,\lambda}(\mathbf{h}) - \mathbf{o}_{j,\beta}(\mathbf{h})\|_2 = O\left(\frac{|(\mathbf{W}_Q \mathbf{h})^\top \mathbf{W}_K \mathbf{r}|}{\sqrt{d_k}}\right).$$

Appendix D.7 gives the proof and explicit bound. Intuitively, stronger steering at one position compensates for the attention received by multiple steered positions, exactly matching their contribution when attention weights remain unchanged. In practice, repeated steering also changes subsequent input representations, going beyond this fixed-state comparison. Nevertheless, Figure 4(b) shows the lowest mean attention-output discrepancy at the predicted strength β^* among those tested. This comparison illustrates strength compensation, but matching full steering may also reproduce its capability effects. We therefore select values from $\{0.01, 0.1, 1, 10\}$ for additive steering and target angles from $\{10^\circ, 30^\circ, 50^\circ, 70^\circ, 90^\circ\}$ for COAST based on the control-capability trade-off.




Figure 5 Control-capability trade-offs across four models and five tasks. Using Additive steering and COAST, we compare one- and five-token Prefix Steering with full steering, prompting, and unsteered baselines. Prefix Steering generally achieves effective control while better preserving general capabilities than full steering. Vertical axes use piecewise scaling and omit empty intervals.

5.2 Evaluation Across Models, Tasks, and Steering Methods

We next evaluate short steering spans across models, tasks, and intervention methods, including final-answer formatting after reasoning, where control is required well after steering ends. Using additive steering [44] and COAST [24], we compare Prefix Steering with full steering, prompting (Appendix C.2), and unsteered baselines on Qwen3 1.7B and 14B [54] and OLMo 3 7B and 32B [25]. Tasks cover safety [21], sentiment [38], politeness [11], and two instruction-following tasks on MATH-500 [16]: IF-Boxed requires the final answer in `\boxed{}`, whereas IF-Plain requires it in plain text. Capability is measured by MMLU-Pro accuracy [48] for concept control and MATH-500 accuracy for formatting.

Figure 5 shows that one- and five-token prefixes often retain much of full steering’s behavioral control while better preserving general capabilities. Compared with prompting, the benefits are task-dependent: prompting performs well on sentiment, politeness, and boxed-answer formatting, whereas steering achieves stronger control on safety and plain-answer formatting in our evaluation, highlighting its value when prompting alone does not reliably elicit the target behavior. Its effectiveness on final-answer formatting after reasoning further shows that a brief intervention can influence behavior required well after steering ends.

Comparison with other strength policies.

The preceding results suggest that limiting intervention can improve the control–capability trade-off. We next ask whether steering should stop after a short prefix or continue with reduced or adaptive strength. On OLMo 3 7B, we compare Prefix Steering with constant full steering, linear and exponential decay, DAS [34], and ACT [47]. Appendix B.3 provides implementation details.

Figure 6 shows competitive trade-offs for short prefixes across the evaluated tasks. Prefix Steering uses a fixed cutoff to emphasize early intervention, consistent with the sharp decrease in DAS strengths after the first few generated tokens observed by Scalena et al. [34]. In our evaluation, longer or adaptive interventions can achieve stronger control, but Prefix Steering offers a more favorable control–capability Pareto frontier.




Figure 6 Steering policies on OLMo 3 7B. Capability is MMLU-Pro accuracy for concept control and MATH-500 accuracy for formatting; control is target-behavior success. Both are percentages.

💡 Takeaway

Short steering spans, sometimes just one token, often retain much of full steering’s behavioral control while better preserving general capabilities. Duration and strength jointly shape this trade-off, with short prefixes remaining competitive with prompting and alternative strength policies across the evaluated settings.

6 Conclusion

We connected prompting and activation steering through conditions for matching attention outputs under fixed-state assumptions. Across models and tasks, Prefix Steering with existing directions often retains much of full steering’s behavioral control while better preserving general capabilities. Its effectiveness on final-answer formatting after reasoning shows that control can outlast intervention. Duration and strength thus jointly shape the control–capability trade-off.

AI use statement

We used generative AI tools to implement and edit research code, refine mathematical statements, and assist with mathematical derivations and proof writing. We also used these tools to draft and revise portions of the manuscript, improve readability, identify relevant literature, and edit LaTeX. The authors independently reviewed and verified the AI-assisted text, code, mathematical claims, and proofs, and take responsibility for all final content and artifacts.

Ethics statement

Our work studies activation steering as a tool for bidirectional behavioral control, allowing target behaviors to be strengthened or suppressed. Such methods are dual use: they may help suppress undesirable behavior, but may also be used to strengthen harmful behavior or produce unintended changes in other capabilities. Our experiments evaluate these effects in controlled benchmark settings and do not establish that a steered model is safe or suitable for deployment. Applying these interventions therefore requires careful evaluation of both behavioral control and capability changes.

Reproducibility statement

The assumptions and complete proofs of our theoretical results are provided in [Appendix D](#) and the subsequent theoretical appendices. [Appendix A](#) specifies the datasets and splits, direction construction, intervention-layer and strength selection, generation settings, and evaluation protocols. The intervention operators, strength policies, and evaluation prompts are documented in [Appendix B](#) and [Appendix C](#). We plan to publicly release our code and reproduction instructions.

References

- [1] Laziz U Abdullaev, Noelle YL Wong, Ryan TZ Lee, Shiqi Jiang, Khoi NM Nguyen, and Tan M Nguyen. Concept heterogeneity-aware representation steering. *arXiv preprint arXiv:2603.02237*, 2026.
- [2] Andy Ardit, Oscar Obeso, Aaquib Syed, Daniel Paleka, Nina Panickssery, Wes Gurnee, and Neel Nanda. Refusal in language models is mediated by a single direction. *Advances in Neural Information Processing Systems*, 37:136037–136083, 2024.
- [3] Yuntao Bai, Saurav Kadavath, Sandipan Kundu, Amanda Askell, Jackson Kernion, Andy Jones, Anna Chen, Anna Goldie, Azalia Mirhoseini, Cameron McKinnon, et al. Constitutional ai: Harmlessness from ai feedback. *arXiv preprint arXiv:2212.08073*, 2022.
- [4] Yuntai Bao, Qinfeng Li, Xinyan Yu, Ge Su, Wenqi Zhang, Liu Yan, Haiqin Weng, Jianwei Yin, and Xuhong Zhang. Towards steering without sacrifice: Principled training of steering vectors for prompt-only interventions. *arXiv preprint arXiv:2605.05983*, 2026.
- [5] Sumanth Dathathri, Andrea Madotto, Janice Lan, Jane Hung, Eric Frank, Piero Molino, Jason Yosinski, and Rosanne Liu. Plug and play language models: A simple approach to controlled text generation. *arXiv preprint arXiv:1912.02164*, 2019.
- [6] Nelson Elhage, Neel Nanda, Catherine Olsson, Tom Henighan, Nicholas Joseph, Ben Mann, Amanda Askell, Yuntao Bai, Anna Chen, Tom Conerly, Nova DasSarma, Dawn Drain, Deep Ganguli, Zac Hatfield-Dodds, Danny Hernandez, Andy Jones, Jackson Kernion, Liane Lovitt, Kamal Ndousse, Dario Amodei, Tom Brown, Jack Clark, Jared Kaplan, Sam McCandlish, and Chris Olah. A mathematical framework for transformer circuits. *Transformer Circuits Thread*, 2021. <https://transformer-circuits.pub/2021/framework/index.html>.
- [7] Jonas Geiping, Sean McLeish, Neel Jain, John Kirchenbauer, Siddharth Singh, Brian Bartoldson, Bhavya Kailkhura, Abhinav Bhatele, and Tom Goldstein. Scaling up test-time compute with latent reasoning: A recurrent depth approach. *Advances in Neural Information Processing Systems*, 38:41340–41391, 2026.
- [8] Shibo Hao, Sainbayar Sukhbaatar, DiJia Su, Xian Li, Zhiting Hu, Jason Weston, and Yuandong Tian. Training large language models to reason in a continuous latent space. *arXiv preprint arXiv:2412.06769*, 2024.
- [9] Roei Hendel, Mor Geva, and Amir Globerson. In-context learning creates task vectors. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 9318–9333, 2023.
- [10] Geert Heyman and Frederik Vandeputte. Steer like the llm: Activation steering that mimics prompting. *arXiv preprint arXiv:2605.03907*, 2026.
- [11] Intel. Politeguard. Hugging Face dataset, 2025. Available at <https://huggingface.co/datasets/Intel/polite-guard>.
- [12] Diancheng Kang, Zheyuan Liu, Ningshan Ma, Yue Huang, Zhaoxuan Tan, and Meng Jiang. Prompt-activation duality: Improving activation steering via attention-level interventions. *arXiv preprint arXiv:2605.10664*, 2026.
- [13] Brian Lester, Rami Al-Rfou, and Noah Constant. The power of scale for parameter-efficient prompt tuning. In *Proceedings of the 2021 conference on empirical methods in natural language processing*, pages 3045–3059, 2021.
- [14] Kenneth Li, Oam Patel, Fernanda Viégas, Hanspeter Pfister, and Martin Wattenberg. Inference-time intervention: Eliciting truthful answers from a language model. *Advances in neural information processing systems*, 36:41451–41530, 2023.
- [15] Xiang Lisa Li and Percy Liang. Prefix-tuning: Optimizing continuous prompts for generation. In *Proceedings of the 59th annual meeting of the association for computational linguistics and the 11th international joint conference on natural language processing (volume 1: Long papers)*, pages 4582–4597, 2021.
- [16] Hunter Lightman, Vineet Kosaraju, Yuri Burda, Harrison Edwards, Bowen Baker, Teddy Lee, Jan Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. Let’s verify step by step. In *International Conference on Learning Representations*, volume 2024, pages 39578–39601, 2024.
- [17] Peter J Liu, Mohammad Saleh, Etienne Pot, Ben Goodrich, Ryan Sepassi, Lukasz Kaiser, and Noam Shazeer. Generating wikipedia by summarizing long sequences. *arXiv preprint arXiv:1801.10198*, 2018.
- [18] Sheng Liu, Haotian Ye, Lei Xing, and James Zou. In-context vectors: Making in context learning more effective and controllable through latent space steering. *arXiv preprint arXiv:2311.06668*, 2023.
- [19] Xiao Liu, Kaixuan Ji, Yicheng Fu, Weng Tam, Zhengxiao Du, Zhilin Yang, and Jie Tang. P-tuning: Prompt tuning can be comparable to fine-tuning across scales and tasks. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, pages 61–68, 2022.
- [20] Iuri Macocco, Pau Rodríguez, Arno Blaas, Luca Zappella, Marco Baroni, and Xavier Suau. On the effectiveness-fluency trade-off in llm conditioning: A systematic study. *arXiv preprint arXiv:2606.12234*, 2026.
- [21] Mantas Mazeika, Long Phan, Xuwang Yin, Andy Zou, Zifan Wang, Norman Mu, Elham Sakhaee, Nathaniel Li, Steven Basart, Bo Li, et al. Harmbench: A standardized evaluation framework for automated red teaming and robust refusal. *arXiv preprint arXiv:2402.04249*, 2024.

- [22] Sewon Min, Xinxu Lyu, Ari Holtzman, Mikel Artetxe, Mike Lewis, Hannaneh Hajishirzi, and Luke Zettlemoyer. Rethinking the role of demonstrations: What makes in-context learning work? In *Proceedings of the 2022 conference on empirical methods in natural language processing*, pages 11048–11064, 2022.
- [23] Dung Viet Nguyen, Yen Pham, Hieu Vu, Lei Zhang, and Tan Nguyen. Activation steering with a feedback controller. In *International Conference on Learning Representations*, volume 2026, pages 154471–154505, 2026.
- [24] Tam Nguyen, Tu Anh Nguyen, Sina Alemohammad, and Richard G Baraniuk. Minimizing collateral damage in activation steering. *arXiv preprint arXiv:2605.01167*, 2026.
- [25] Team Olmo, Allyson Ettinger, Amanda Bertsch, Bailey Kuehl, David Graham, David Heineman, Dirk Groeneveld, Faeze Brahman, Finbarr Timbers, Hamish Ivison, et al. Olmo 3. *arXiv preprint arXiv:2512.13961*, 2025.
- [26] Catherine Olsson, Nelson Elhage, Neel Nanda, Nicholas Joseph, Nova DasSarma, Tom Henighan, Ben Mann, Amanda Askell, Yuntao Bai, Anna Chen, et al. In-context learning and induction heads. *arXiv preprint arXiv:2209.11895*, 2022.
- [27] Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. Training language models to follow instructions with human feedback. *Advances in neural information processing systems*, 35:27730–27744, 2022.
- [28] Aleksandar Petrov, Philip Torr, and Adel Bibi. When do prompting and prefix-tuning work? a theory of capabilities and limitations. In *International Conference on Learning Representations*, volume 2024, pages 6031–6054, 2024.
- [29] Joris Postmus and Steven Abreu. Steering large language models using conceptors: Improving addition-based activation engineering. *arXiv preprint arXiv:2410.16314*, 2024.
- [30] Laria Reynolds and Kyle McDonell. Prompt programming for large language models: Beyond the few-shot paradigm. In *Extended abstracts of the 2021 CHI conference on human factors in computing systems*, pages 1–7, 2021.
- [31] Nina Rimskey, Nick Gabrieli, Julian Schulz, Meg Tong, Evan Hubinger, and Alexander Turner. Steering llama 2 via contrastive activation addition. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 15504–15522, 2024.
- [32] Pau Rodriguez, Arno Blaas, Michal Klein, Luca Zappella, Nicholas Apostoloff, Xavier Suau, et al. Controlling language and diffusion models by transporting activations. In *International Conference on Learning Representations*, volume 2025, pages 89812–89855, 2025.
- [33] Victor Sanh, Albert Webson, Colin Raffel, Stephen H Bach, Lintang Sutawika, Zaid Alyafeai, Antoine Chaffin, Arnaud Stiegler, Teven Le Scao, Arun Raja, et al. Multitask prompted training enables zero-shot task generalization. *arXiv preprint arXiv:2110.08207*, 2021.
- [34] Daniel Scalena, Gabriele Sarti, and Malvina Nissim. Multi-property steering of large language models with dynamic activation composition. In *Proceedings of the 7th BlackboxNLP Workshop: Analyzing and Interpreting Neural Networks for NLP*, pages 577–603, 2024.
- [35] Sander Schulhoff, Michael Ilie, Nishant Balepur, Konstantine Kahadze, Amanda Liu, Chenglei Si, Yinheng Li, Aayush Gupta, Hyojung Han, Sevien Schulhoff, et al. The prompt report: a systematic survey of prompt engineering techniques. *arXiv preprint arXiv:2406.06608*, 2024.
- [36] Abhay Sheshadri, Aidan Ewart, Phillip Guo, Aengus Lynch, Cindy Wu, Vivek Hebbar, Henry Sleight, Asa Cooper Stickland, Ethan Perez, Dylan Hadfield-Menell, et al. Latent adversarial training improves robustness to persistent harmful behaviors in llms. *arXiv preprint arXiv:2407.15549*, 2024.
- [37] Taylor Shin, Yasaman Razeghi, Robert L Logan Iv, Eric Wallace, and Sameer Singh. Autoprompt: Eliciting knowledge from language models with automatically generated prompts. In *Proceedings of the 2020 conference on empirical methods in natural language processing (EMNLP)*, pages 4222–4235, 2020.
- [38] Richard Socher, Alex Perelygin, Jean Wu, Jason Chuang, Christopher D Manning, Andrew Y Ng, and Christopher Potts. Recursive deep models for semantic compositionality over a sentiment treebank. In *Proceedings of the 2013 conference on empirical methods in natural language processing*, pages 1631–1642, 2013.
- [39] Alessandro Stolfo, Vidhisha Balachandran, Safoora Yousefi, Eric Horvitz, and Besmira Nushi. Improving instruction-following in language models through activation steering. In *International Conference on Learning Representations*, volume 2025, pages 55790–55823, 2025.
- [40] Jianlin Su, Murtadha Ahmed, Yu Lu, Shengfeng Pan, Wen Bo, and Yunfeng Liu. Roformer: Enhanced transformer with rotary position embedding. *Neurocomputing*, 568:127063, 2024.
- [41] Daniel Tan, David Chanin, Aengus Lynch, Brooks Paige, Dimitrios Kanoulas, Adrià Garriga-Alonso, and Robert Kirk. Analysing the generalisation and reliability of steering vectors. *Advances in Neural Information Processing Systems*, 37:139179–139212, 2024.

- [42] Eric Todd, Millicent Li, Arnab Sen Sharma, Aaron Mueller, Byron Wallace, and David Bau. Function vectors in large language models. In *International conference on learning representations*, volume 2024, pages 17282–17333, 2024.
- [43] Brown Tom, Mann Benjamin, Ryder Nick, Subbiah Melanie, Dhariwal Prafulla, Neelakantan Arvind, Shyam Pranav, Sastry Girish, Askell Amanda, Agarwal Sandhini, et al. Language models are few-shot learners. *Advances in neural information processing systems*, 33(1):1877–1901, 2020.
- [44] Alexander Matt Turner, Lisa Thiergart, Gavin Leech, David Udell, Juan J Vazquez, Ulisse Mini, and Monte MacDiarmid. Steering language models with activation engineering. *arXiv preprint arXiv:2308.10248*, 2023.
- [45] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. *Advances in neural information processing systems*, 30, 2017.
- [46] Lean Wang, Lei Li, Damai Dai, Deli Chen, Hao Zhou, Fandong Meng, Jie Zhou, and Xu Sun. Label words are anchors: An information flow perspective for understanding in-context learning. In Houda Bouamor, Juan Pino, and Kalika Bali, editors, *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 9840–9855, Singapore, December 2023. Association for Computational Linguistics. doi: 10.18653/v1/2023.emnlp-main.609. URL <https://aclanthology.org/2023.emnlp-main.609/>.
- [47] Tianlong Wang, Xianfeng Jiao, Yinghao Zhu, Zhongzhi Chen, Yifan He, Xu Chu, Junyi Gao, Yasha Wang, and Liantao Ma. Adaptive activation steering: A tuning-free llm truthfulness improvement method for diverse hallucinations categories. In *Proceedings of the ACM on Web Conference 2025*, pages 2562–2578, 2025.
- [48] Yubo Wang, Xueguang Ma, Ge Zhang, Yuansheng Ni, Abhranil Chandra, Shiguang Guo, Weiming Ren, Aaran Arulraj, Xuan He, Ziyang Jiang, et al. Mmlu-pro: A more robust and challenging multi-task language understanding benchmark. *Advances in Neural Information Processing Systems*, 37:95266–95290, 2024.
- [49] Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Fei Xia, Ed Chi, Quoc V Le, Denny Zhou, et al. Chain-of-thought prompting elicits reasoning in large language models. *Advances in neural information processing systems*, 35:24824–24837, 2022.
- [50] Zhengxuan Wu, Aryaman Arora, Zheng Wang, Atticus Geiger, Dan Jurafsky, Christopher D Manning, and Christopher Potts. Reft: Representation finetuning for language models. *Advances in Neural Information Processing Systems*, 37:63908–63962, 2024.
- [51] Zhengxuan Wu, Atticus Geiger, Aryaman Arora, Jing Huang, Zheng Wang, Noah Goodman, Christopher D Manning, and Christopher Potts. pyvene: A library for understanding and improving pytorch models via interventions. In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 3: System Demonstrations)*, pages 158–165, 2024.
- [52] Zhengxuan Wu, Aryaman Arora, Atticus Geiger, Zheng Wang, Jing Huang, Dan Jurafsky, Christopher D Manning, and Christopher Potts. Axbench: Steering llms? even simple baselines outperform sparse autoencoders. *arXiv preprint arXiv:2501.17148*, 2025.
- [53] Zhengxuan Wu, Qinan Yu, Aryaman Arora, Christopher D Manning, and Chris Potts. Improved representation steering for language models. *Advances in Neural Information Processing Systems*, 38:160589–160641, 2026.
- [54] An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Gao, Chengen Huang, Chenxu Lv, et al. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.
- [55] ZeJia You, Chunyuan Deng, and Hanjie Chen. Spherical steering: Geometry-aware activation rotation for language models. *arXiv preprint arXiv:2602.08169*, 2026.
- [56] Biao Zhang and Rico Sennrich. Root mean square layer normalization. *Advances in neural information processing systems*, 32, 2019.
- [57] Hanqing Zhang, Haolin Song, Shaoyu Li, Ming Zhou, and Dawei Song. A survey of controllable text generation using transformer-based pre-trained language models. *ACM Computing Surveys*, 56(3):1–37, 2023.
- [58] Andy Zou, Long Phan, Sarah Chen, James Campbell, Phillip Guo, Richard Ren, Alexander Pan, Xuwang Yin, Mantas Mazeika, Ann-Kathrin Dombrowski, et al. Representation engineering: A top-down approach to ai transparency. *arXiv preprint arXiv:2310.01405*, 2023.

Appendix

Contents

A	Experimental Details	16
B	Steering Method Details	17
B.1	From Constructed Displacements to DiM Directions	17
B.2	Intervention Operators	18
B.3	Steering Strength Policies	18
B.4	Direction Estimation with DiM	19
C	Evaluation and Baseline Prompts	19
C.1	LLM-as-a-Judge Prompts	19
C.2	Prompting Baseline	20
D	Proofs of the Main Results	20
D.1	Proof of Theorem 1	20
D.2	Proof of Lemma 1	21
D.3	Proof of Lemma 2	23
D.4	Proof of Theorem 2	24
D.5	Proof of Lemma 3	25
D.6	Proof of Lemma 4	26
D.7	Proof of Lemma 5	27
E	Multi-Head Attention-Output Matching	29
E.1	Multi-head output matching	29
E.2	Examples illustrating multi-head matching	31

The appendix provides experimental details in [Appendix A](#), steering methods in [Appendix B](#), and evaluation and baseline prompts in [Appendix C](#). [Appendix D](#) gives the proofs of the main results, and [Appendix E](#) extends the attention analysis to multiple heads.

A Experimental Details

Concept-control datasets. Direction construction and evaluation use disjoint data. For each concept, we construct a DiM direction from 100 positive and 100 negative training sentences, using the residual-stream activation at the final prompt token.

- **Safety.** We construct safety directions using LLM-LAT/*benign-dataset* as the positive set and LLM-LAT/*harmful-dataset* as the negative set [36], and evaluate safety behavior on Harm-Bench [21].
- **Sentiment.** We construct sentiment directions using the SST-2 training split and evaluate on its validation split [38].
- **Politeness.** We construct politeness directions using the PoliteGuard training split and evaluate on its test split [11]. We retain only the *polite* and *impolite* labels.

Instruction-following tasks. We use MATH-500 [16] to test whether interventions at the beginning of generation can control final-answer formats after reasoning: “The answer is \boxed{ }” and “The answer is . . .”. Both tasks share a partition of 50 problems for direction construction, 50 for intervention layer and strength selection, and 400 for evaluation, with no overlap. For each format and candidate layer, we construct a difference-in-means (DiM) direction [31] from final prompt token activations,

subtracting the mean for problems without the format instruction from the mean for the same problems with it appended.

Steering duration. Prefix Steering starts at the final prompt token and ends after the specified number of steered positions; full steering continues throughout generation. Single-token Prefix Steering modifies only the final prompt token.

Intervention layer and strength. We reserve 50 test examples as a shared tuning set and exclude them from final evaluation. For each model–concept pair and intervention operator, we jointly select the intervention layer and strength on this tuning set. Within each operator, full steering and Prefix Steering use the same layer. We consider nine candidate layers at 10%, 20%, ..., 90% of the model’s depth. We normalize each direction to unit ℓ_2 norm. For additive steering, we use a coefficient $c = s\bar{r}$, where $s \in \{0.01, 0.1, 1, 10\}$ and $\bar{r}$ is the mean ℓ_2 norm of residual-stream activations at the candidate layer over the positive and negative samples used for DiM construction. For COAST, we consider target angles $\theta \in \{10^\circ, 30^\circ, 50^\circ, 70^\circ, 90^\circ\}$ between the steered hidden state and the steering direction, setting the cosine-similarity parameter to $\alpha = \cos \theta$. We select the layer–strength combination maximizing concept score while retaining at least 80% of the unsteered general score.

Generation and evaluation protocol. We use Gemini 3.5 Flash Lite as the LLM judge, with task-specific prompts described in [Appendix C.1](#). We use greedy decoding with limits of 512 generated tokens for concept-control tasks, 1024 for MMLU-Pro, and 4096 for MATH-500, with the same decoding settings across methods within each task. For concept-control tasks, we evaluate only elicitation: safe behavior on HarmBench prompts [21], and positive sentiment and politeness on held-out negative-class inputs. For MATH-500, we evaluate adherence to the target final-answer format. We measure capability using MMLU-Pro multiple-choice accuracy [48] for concept-control tasks and MATH-500 answer accuracy for final-answer formatting.

B Steering Method Details

B.1 From Constructed Displacements to DiM Directions

We compare DiM directions with a displacement $\mathbf{r}$ obtained by jointly solving across all attention heads, with the objective of matching the attention output at the first generated token to that of the prompted run. This comparison examines the connection between our attention-level analysis in [Section 4.1](#) and practical steering. The construction depends on the specific prompt, context, and input representation, so the resulting displacement generally changes with these conditions. Moreover, extending the construction to multiple attention heads requires compatibility conditions that need not hold, as discussed in [Appendix E.1](#), and an exact matching displacement may not exist. DiM instead estimates a concept direction from a collection of examples, yielding a direction that can be reused across inputs for a given model, layer, and concept. We evaluate safety on 100 examples, using the safety prompt in [Appendix C.2](#) appended to the end of each input rather than supplied as a system message. [Figure 7](#) shows positive cosine similarities between DiM and the constructed displacement across the evaluated layers of OLMo 3 7B, indicating partial directional alignment in this setting. This alignment suggests a geometric connection, although it does not imply that DiM satisfies the exact attention-matching conditions. We therefore use DiM directions to examine whether the intuition from our analysis extends to practical steering: can a brief initial intervention retain the behavioral control of full steering while better preserving general capabilities? We leave the development of practical steering directions derived directly from prompts through our construction, and reusable across inputs, to future work.




Figure 7 Cosine similarity between the DiM direction and the constructed displacement $\mathbf{r}$ across relative layer depths in OLMo 3 7B.

B.2 Intervention Operators

We consider two intervention operators: additive steering translates the hidden state along a given direction, while COAST controls its angle to that direction while preserving its norm. Let $\mathbf{r}$ be a unit steering direction. Both operators can be applied over a prefix or throughout generation.

Additive steering [44]. Additive steering adds the direction $\mathbf{r}$ to the hidden state with strength c :

$$\mathbf{h}'_k = \mathbf{h}_k + c \mathbf{r}, \quad c \in \mathbb{R}.$$

COAST [24]. COAST preserves the hidden-state norm and sets its cosine similarity with $\mathbf{r}$ to a target $\alpha \in [-1, 1]$. For $\mathbf{h}_k \neq 0$, let $\hat{\mathbf{h}}_k = \mathbf{h}_k / \|\mathbf{h}_k\|_2$. The update is

$$\mathbf{x}^* \in \arg \min_{\|\mathbf{x}\|_2=1, \mathbf{r}^\top \mathbf{x}=\alpha} (\mathbf{x} - \hat{\mathbf{h}}_k)^\top \Sigma (\mathbf{x} - \hat{\mathbf{h}}_k), \quad \mathbf{h}'_k = \|\mathbf{h}_k\|_2 \mathbf{x}^*,$$

where Σ weights changes to non-target features and is estimated from the second moment of reference activations. Thus $\|\mathbf{h}'_k\|_2 = \|\mathbf{h}_k\|_2$ and $\cos(\mathbf{h}'_k, \mathbf{r}) = \alpha$: COAST controls angular alignment while minimizing the weighted collateral change under these constraints.

B.3 Steering Strength Policies

We describe adaptive strength rules and position-dependent schedules, using c_t for the effective additive coefficient to distinguish it from COAST’s target cosine similarity α .

Dynamic Activation Composition (DAS) [34]. At each generation step t , this method compares the next-token distributions without steering and with a probe intervention of strength $c_{\max} > 0$, a hyperparameter that also caps the final coefficient.² It takes the union of their nucleus token sets and renormalizes both distributions over this shared support, yielding $\tilde{p}_t^0$ and $\tilde{p}_t^{c_{\max}}$. The coefficient for each target property is then

$$c_t = \min\{D_{\text{KL}}(\tilde{p}_t^0 \parallel \tilde{p}_t^{c_{\max}}), c_{\max}\}.$$

Small distributional changes yield weak steering. This rule requires extra forward evaluations and need not produce monotonically decreasing strengths.

Adaptive Activation Steering (ACT) [47]. We adopt ACT’s adaptive coefficient rule with a fixed difference-in-means steering vector $\mathbf{v}$. Let $p_\theta(\mathbf{h}_t) \in [0, 1]$ be a probe’s target-concept score for the current hidden state $\mathbf{h}_t$. The intervention is

$$\mathbf{h}'_t = \mathbf{h}_t + c_t \mathbf{v}, \quad c_t = A(1 - p_\theta(\mathbf{h}_t) + \beta),$$

where A and β are global hyperparameters.³ Lower probe scores yield stronger steering.

Prefix Steering. Let $t = 0$ denote the final input token and $t \geq 1$ the subsequent generated tokens. For additive Prefix Steering, the coefficient is constant over the first k positions and zero afterward:

$$c_t = \begin{cases} c_0, & 0 \leq t < k, \\ 0, & t \geq k. \end{cases}$$

Here $k \geq 1$ is the prefix length; $k = 1$ intervenes only at the final input token. For COAST, we apply the operator with a fixed target cosine similarity over the first k positions and skip the operator afterward; setting the target cosine similarity to zero would still modify the activation.

Linear Decay. The coefficient decreases linearly from c_0 to zero over a horizon $K > 0$:

$$c_t = c_0 \max\left\{1 - \frac{t}{K}, 0\right\}.$$

This schedule gradually reduces the intervention and stops it at $t = K$.

²We use $c_{\max} = 2$, following the original method.

³We use $A = 12$ and $\beta = 0$, following the full-data setting of the original method.

Exponential Decay. The coefficient decays exponentially with time constant $\tau > 0$:

$$c_t = c_0 \exp\left(-\frac{t}{\tau}\right).$$

We set $\tau = 128$. Unlike the finite-horizon schedules above, exponential decay retains a nonzero intervention at every finite step when $c_0 \neq 0$.

B.4 Direction Estimation with DiM

Difference in Means (DiM) [31] estimates a concept direction from positive and negative examples. At a fixed layer and token position, we average the activations within each group to obtain μ_c^+ and μ_c^- . We then subtract the negative mean from the positive mean and normalize the difference:

$$\hat{r}_c = \frac{\mu_c^+ - \mu_c^-}{\|\mu_c^+ - \mu_c^-\|_2}.$$

Under the linear decomposition in the background section, this difference approximately isolates the target direction when the other components have similar means across the two groups.

C Evaluation and Baseline Prompts

C.1 LLM-as-a-Judge Prompts

We use Gemini 3.5 Flash Lite with the five templates below. The judge treats input as data and ignores embedded instructions. Each prompt returns only the target label or OTHERS; empty or unintelligible responses receive OTHERS. The judge is not told the steering method or strength.

✈ Safety judge prompt

Return SAFE if the assistant response refuses the harmful request or gives a benign response without substantive harmful assistance; otherwise return OTHERS. A disclaimer alone does not make a harmful response safe. Return only the label.

Harmful request: {prompt}
Assistant response: {response}

✈ Sentiment judge prompt

Judge the sentiment of the assistant response itself, not the user prompt. Return POSITIVE if the response expresses clearly positive sentiment; otherwise return OTHERS. Return only the label.

User prompt: {prompt}
Assistant response: {response}

✈ Politeness judge prompt

Return POLITE if the assistant response is clearly respectful and courteous toward its addressee; otherwise return OTHERS. Judge tone, not agreement with the user. Return only the label.

User prompt: {prompt}
Assistant response: {response}

✈ Boxed final-answer judge prompt

Return MATCH only if the response ends with a standalone final-answer line of exactly “The answer is \boxed{ANSWER}”, with a nonempty answer inside the box. Earlier reasoning may precede this line, but an occurrence of this phrase within the reasoning does not count, and no additional text may follow the final-answer line. Ignore mathematical correctness. Otherwise return OTHERS. Return only the label.

Assistant response: {response}

↗ Plain final-answer judge prompt

Return MATCH only if the response ends with a standalone final-answer line of the form “The answer is ANSWER”, with a nonempty answer immediately following “The answer is” in plain text. Earlier reasoning may precede this line, but an occurrence of this phrase within the reasoning does not count, and no additional text may follow the final-answer line. Ignore mathematical correctness. Otherwise return OTHERS. Return only the label.

Assistant response: {response}

C.2 Prompting Baseline

For the prompting baseline, we use the following five task-specific system prompts. Each prompt is supplied as the system message, with the task input supplied as the user message.

↗ Safety system prompt

Respond safely. Refuse harmful requests and do not provide harmful assistance.

↗ Sentiment system prompt

Respond with clearly positive sentiment.

↗ Politeness system prompt

Respond politely and respectfully.

↗ IF-Boxed system prompt

Solve the problem step by step and write your final answer as “The answer is \boxed{ANSWER}”, replacing ANSWER with your final answer.

↗ IF-Plain system prompt

Solve the problem step by step and write your final answer as “The answer is ANSWER”, replacing ANSWER with your final answer in plain text.

D Proofs of the Main Results

The proofs follow the order of the statements in Section 4 and Section 5. Throughout, $q = W_Q h$ is the reference query and $q' = W_Q h'$ is a comparison query. Attention outputs are written as functions of h or h' , as in the main text; q and q' are used to compute attention weights and scores.

D.1 Proof of Theorem 1

Notation and setting. We use linear key and value projections, omitting positional effects and pre-projection normalization. Let $q = W_Q h$ denote the query produced by the fixed representation h in Theorem 1; it is shared by both attention calculations, without assuming their actual queries coincide. Set $k_{t_i} = W_K h_{t_i}$, $v_{t_i} = W_V h_{t_i}$, $k_p = W_K h_p$, $v_p = W_V h_p$, and $\Delta h = h_p - h_{t_n}$, $\Delta k = k_p - k_{t_n}$, $\Delta v = v_p - v_{t_n}$. The shared positions are $\mathcal{I} = \{1, \dots, n-1\}$, with identical keys and values in both settings. At the reference query, write $a_{t_n} = e^{q^\top k_{t_n} / \sqrt{d_k}}$ and $a_p = e^{q^\top k_p / \sqrt{d_k}}$. For any group of positions, we use the sum of its unnormalized attention weights and the average of its values weighted by those weights.

Theorem 1. Suppose $(W_Q h)^\top W_K z \neq 0$ for some $z \in \ker(W_V)$. There exists $r = r(h; h_{t_n}, h_p, W_Q, W_K, W_V)$ such that steering h_{t_n} by r gives $o_{\text{steer}}(h) = o_{\text{prompt}}(h)$.

The displacement

$$\mathbf{r} = \frac{a_p}{a_{t_n} + a_p} \Delta \mathbf{h} + \frac{\sqrt{d_k} \log\left(\frac{a_{t_n} + a_p}{a_{t_n}}\right) - \frac{a_p}{a_{t_n} + a_p} \mathbf{q}^\top \mathbf{W}_K \Delta \mathbf{h}}{\mathbf{q}^\top \mathbf{W}_K \mathbf{z}} \mathbf{z} \quad (1)$$

satisfies $\mathbf{o}_{\text{steer}}(\mathbf{h}) = \mathbf{o}_{\text{prompt}}(\mathbf{h})$.

Proof. At the fixed query, the prompted positions have weights

$$a_{t_n} = \exp\left(\frac{\mathbf{q}^\top \mathbf{W}_K \mathbf{h}_{t_n}}{\sqrt{d_k}}\right), \quad a_p = \exp\left(\frac{\mathbf{q}^\top \mathbf{W}_K \mathbf{h}_p}{\sqrt{d_k}}\right). \quad (2)$$

The steered position has key $\mathbf{W}_K(\mathbf{h}_{t_n} + \mathbf{r})$, value $\mathbf{W}_V(\mathbf{h}_{t_n} + \mathbf{r})$, and weight

$$a'_{t_n} = a_{t_n} \exp\left(\frac{\mathbf{q}^\top \mathbf{W}_K \mathbf{r}}{\sqrt{d_k}}\right). \quad (3)$$

To obtain equal outputs, we seek a displacement satisfying two sufficient conditions: $a'_{t_n} = a_{t_n} + a_p$ and $\mathbf{v}'_{t_n} = (a_{t_n} \mathbf{v}_{t_n} + a_p \mathbf{v}_p) / (a_{t_n} + a_p)$. They make the steered position contribute the same numerator and denominator as the two prompted positions.

First make the steered value equal to the attention-weighted average of the two original values. Since $\mathbf{v}_p - \mathbf{v}_{t_n} = \mathbf{W}_V \Delta \mathbf{h}$ and $\mathbf{v}'_{t_n} = \mathbf{v}_{t_n} + \mathbf{W}_V \mathbf{r}$, this requires

$$\mathbf{W}_V \mathbf{r} = \frac{a_p}{a_{t_n} + a_p} \mathbf{W}_V \Delta \mathbf{h}. \quad (4)$$

Every solution of this equation has the form

$$\mathbf{r} = \frac{a_p}{a_{t_n} + a_p} \Delta \mathbf{h} + \mathbf{u}, \quad \mathbf{u} \in \ker(\mathbf{W}_V). \quad (5)$$

Such a $\mathbf{u}$ can change the attention score without changing this value.

Next, matching the weights requires $a'_{t_n} = a_{t_n} + a_p$. Taking logarithms gives

$$\mathbf{q}^\top \mathbf{W}_K \mathbf{r} = \sqrt{d_k} \log\left(\frac{a_{t_n} + a_p}{a_{t_n}}\right). \quad (6)$$

Substituting (5) shows precisely what $\mathbf{u}$ must supply:

$$\mathbf{q}^\top \mathbf{W}_K \mathbf{u} = \sqrt{d_k} \log\left(\frac{a_{t_n} + a_p}{a_{t_n}}\right) - \frac{a_p}{a_{t_n} + a_p} \mathbf{q}^\top \mathbf{W}_K \Delta \mathbf{h}. \quad (7)$$

By assumption, $\mathbf{z} \in \ker(\mathbf{W}_V)$ and $\mathbf{q}^\top \mathbf{W}_K \mathbf{z} \neq 0$, so a suitable multiple of $\mathbf{z}$ supplies the required score change without changing the value. Substituting this choice of $\mathbf{u}$ into (5) gives the displacement in (1), for which

$$a'_{t_n} = a_{t_n} + a_p, \quad a'_{t_n} \mathbf{v}'_{t_n} = a_{t_n} \mathbf{v}_{t_n} + a_p \mathbf{v}_p.$$

Thus, the steered token contributes exactly as much to the attention denominator and numerator as t_n and p together. All other positions are unchanged, so $\mathbf{o}_{\text{steer}}(\mathbf{h}) = \mathbf{o}_{\text{prompt}}(\mathbf{h})$. $\square$

Scope across layers. Earlier interventions can change the query and context representations at later layers. The theorem holds these states fixed, so it does not establish matching across a full forward pass.

D.2 Proof of Lemma 1

Lemma 1. Fix the context representations and the displacement $\mathbf{r}$ constructed in Theorem 1. Define

$$\mathcal{U} = \text{span}\{\mathbf{W}_K(\mathbf{h}_p - \mathbf{h}_{t_n}), \mathbf{W}_K \mathbf{r}\}$$

with $\dim \mathcal{U}^\perp \in \{d_k - 2, d_k - 1\}$. Then any representation $\mathbf{h}' \in \mathbb{R}^d$ satisfying $\mathbf{W}_Q(\mathbf{h}' - \mathbf{h}) \in \mathcal{U}^\perp$ preserves the match:

$$\mathbf{o}_{\text{steer}}(\mathbf{h}') = \mathbf{o}_{\text{prompt}}(\mathbf{h}').$$

Proof. Set $\mathbf{q} = \mathbf{W}_Q \mathbf{h}$, $\mathbf{q}' = \mathbf{W}_Q \mathbf{h}'$, and $\delta \mathbf{q} = \mathbf{q}' - \mathbf{q}$. We first explain the choice of $\mathcal{U}$. Let a'_{t_n} denote the unnormalized attention weight after steering t_n . At the reference query, the construction gives

$$\mathbf{v}_{t_n} + \mathbf{W}_V \mathbf{r} = \frac{a_{t_n}(\mathbf{q})\mathbf{v}_{t_n} + a_p(\mathbf{q})\mathbf{v}_p}{a_{t_n}(\mathbf{q}) + a_p(\mathbf{q})}, \quad a'_{t_n}(\mathbf{q}) = a_{t_n}(\mathbf{q}) + a_p(\mathbf{q}). \quad (8)$$

The first equality is preserved if the relative weight a_p/a_{t_n} stays fixed: the two original values then retain the same weighted average. Dividing the second equality by a_{t_n} gives $a'_{t_n}/a_{t_n} = 1 + a_p/a_{t_n}$. Once the first ratio is fixed, preserving a'_{t_n}/a_{t_n} therefore preserves the second equality as well.

These two ratios depend on the query through two key directions:

$$\log \frac{a_p(\mathbf{q}')}{a_{t_n}(\mathbf{q}')} = \frac{\mathbf{q}'^\top \Delta \mathbf{k}}{\sqrt{d_k}}, \quad \log \frac{a'_{t_n}(\mathbf{q}')}{a_{t_n}(\mathbf{q}')} = \frac{\mathbf{q}'^\top \mathbf{W}_K \mathbf{r}}{\sqrt{d_k}},$$

where $\Delta \mathbf{k} = \mathbf{k}_p - \mathbf{k}_{t_n} = \mathbf{W}_K(\mathbf{h}_p - \mathbf{h}_{t_n})$. The direction $\Delta \mathbf{k}$ controls the relative weight of the two original tokens; $\mathbf{W}_K \mathbf{r}$ controls the weight multiplier introduced by steering. Thus both ratios stay fixed when

$$\delta \mathbf{q}^\top \Delta \mathbf{k} = 0, \quad \delta \mathbf{q}^\top \mathbf{W}_K \mathbf{r} = 0.$$

This explains the definition

$$\mathcal{U} = \text{span}\{\Delta \mathbf{k}, \mathbf{W}_K \mathbf{r}\} = \text{span}\{\mathbf{W}_K(\mathbf{h}_p - \mathbf{h}_{t_n}), \mathbf{W}_K \mathbf{r}\}.$$

Being orthogonal to $\mathcal{U}$ is exactly being orthogonal to both directions above. Any $\mathbf{h}'$ with $\mathbf{W}_Q(\mathbf{h}' - \mathbf{h}) \in \mathcal{U}^\perp$ therefore gives an allowed query $\mathbf{q}' = \mathbf{W}_Q \mathbf{h}'$ in $\mathbf{q} + \mathcal{U}^\perp$. This is a sufficient condition for output matching. The following two steps verify it explicitly.

Step 1: preserve the weighted average value. The relative weight of p and t_n depends only on their key difference. The first orthogonality condition gives

$$\frac{a_p(\mathbf{q}')}{a_{t_n}(\mathbf{q}')} = e^{\mathbf{q}'^\top \Delta \mathbf{k} / \sqrt{d_k}} = e^{\mathbf{q}^\top \Delta \mathbf{k} / \sqrt{d_k}} = \frac{a_p(\mathbf{q})}{a_{t_n}(\mathbf{q})}.$$

Thus the two tokens retain the same relative weights, even though their individual weights may change. Their values are fixed, so their weighted average also stays the same:

$$\frac{a_{t_n}(\mathbf{q}')\mathbf{v}_{t_n} + a_p(\mathbf{q}')\mathbf{v}_p}{a_{t_n}(\mathbf{q}') + a_p(\mathbf{q}')} = \frac{a_{t_n}(\mathbf{q})\mathbf{v}_{t_n} + a_p(\mathbf{q})\mathbf{v}_p}{a_{t_n}(\mathbf{q}) + a_p(\mathbf{q})} = \mathbf{v}_{t_n} + \mathbf{W}_V \mathbf{r}.$$

Step 2: preserve the sum of weights. Steering adds $\mathbf{W}_K \mathbf{r}$ to the key of t_n . Its weight is therefore multiplied by $e^{\mathbf{q}'^\top \mathbf{W}_K \mathbf{r} / \sqrt{d_k}}$. The second orthogonality condition keeps this multiplier unchanged:

$$\frac{a'_{t_n}(\mathbf{q}')}{a_{t_n}(\mathbf{q}')} = e^{\mathbf{q}'^\top \mathbf{W}_K \mathbf{r} / \sqrt{d_k}} = e^{\mathbf{q}^\top \mathbf{W}_K \mathbf{r} / \sqrt{d_k}} = \frac{a'_{t_n}(\mathbf{q})}{a_{t_n}(\mathbf{q})}.$$

At $\mathbf{q}$, this ratio equals $1 + a_p(\mathbf{q})/a_{t_n}(\mathbf{q})$ by (8). Step 1 preserves the ratio on the right, so

$$\frac{a'_{t_n}(\mathbf{q}')}{a_{t_n}(\mathbf{q}')} = 1 + \frac{a_p(\mathbf{q}')}{a_{t_n}(\mathbf{q}')}, \quad a'_{t_n}(\mathbf{q}') = a_{t_n}(\mathbf{q}') + a_p(\mathbf{q}').$$

Together, Steps 1 and 2 also give

$$a'_{t_n}(\mathbf{q}')(\mathbf{v}_{t_n} + \mathbf{W}_V \mathbf{r}) = a_{t_n}(\mathbf{q}')\mathbf{v}_{t_n} + a_p(\mathbf{q}')\mathbf{v}_p.$$

Hence the steered token contributes the same denominator and numerator as the pair at $\mathbf{q}'$. All other positions are identical between the two calculations at this query, so

$$\mathbf{o}_{\text{steer}}(\mathbf{h}') = \mathbf{o}_{\text{prompt}}(\mathbf{h}') \quad \text{whenever } \mathbf{W}_Q(\mathbf{h}' - \mathbf{h}) \in \mathcal{U}^\perp.$$

Finally, the reference construction gives $\mathbf{q}^\top \mathbf{W}_K \mathbf{r} / \sqrt{d_k} = \log((a_{t_n}(\mathbf{q}) + a_p(\mathbf{q}))/a_{t_n}(\mathbf{q})) > 0$. Thus $\mathbf{W}_K \mathbf{r} \neq 0$. Since $\mathcal{U}$ is spanned by two vectors and contains the nonzero vector $\mathbf{W}_K \mathbf{r}$, its dimension is one or two. Subtracting from d_k gives $\dim \mathcal{U}^\perp \in \{d_k - 2, d_k - 1\}$. $\square$

D.3 Proof of Lemma 2

Lemma 2. Under the same assumptions as Lemma 1, any representation $\mathbf{h}' \in \mathbb{R}^d$ satisfies

$$\|\mathbf{o}_{\text{steer}}(\mathbf{h}') - \mathbf{o}_{\text{prompt}}(\mathbf{h}')\|_2 = O(\text{dist}(\mathbf{W}_Q \mathbf{h}', \mathbf{W}_Q \mathbf{h} + \mathcal{U}^\perp)),$$

where $\mathcal{U} = \text{span}\{\mathbf{W}_K(\mathbf{h}_p - \mathbf{h}_{t_n}), \mathbf{W}_K \mathbf{r}\}$ and dist is Euclidean distance to a set.

Explicit bound. Keep the same queries, values, and displacement as in Appendix D.2. Recall $\Delta \mathbf{k} = \mathbf{k}_p - \mathbf{k}_{t_n}$ and $\Delta \mathbf{v} = \mathbf{v}_p - \mathbf{v}_{t_n}$. The two changes in relative attention scores are

$$\epsilon_k = \frac{|(\mathbf{q}' - \mathbf{q})^\top \Delta \mathbf{k}|}{\sqrt{d_k}}, \quad \epsilon_r = \frac{|(\mathbf{q}' - \mathbf{q})^\top \mathbf{W}_K \mathbf{r}|}{\sqrt{d_k}}.$$

Let d_{val} be the largest Euclidean distance between any two original values:

$$d_{\text{val}} = \max_{i,j \in \{t_1, \dots, t_n, p\}} \|\mathbf{v}_i - \mathbf{v}_j\|_2.$$

This measures the spread of the values, not the query's distance to the matching set. Let w_s be the fraction of attention assigned to the steered token at $\mathbf{q}'$:

$$w_s = \frac{a'_{t_n}(\mathbf{q}')}{a'_{t_n}(\mathbf{q}') + \sum_{i \in \mathcal{I}} a_{t_i}(\mathbf{q}')}.$$

Then the following bound holds, including when $\mathcal{I}$ is empty:

$$\|\mathbf{o}_{\text{steer}}(\mathbf{h}') - \mathbf{o}_{\text{prompt}}(\mathbf{h}')\|_2 \leq \min \left\{ d_{\text{val}}, \frac{w_s \|\Delta \mathbf{v}\|_2 + d_{\text{val}}}{4} \epsilon_k + \frac{d_{\text{val}}}{4} \epsilon_r \right\}. \quad (9)$$

Proof. The prompted pair's weighted average value is

$$\bar{\mathbf{v}}_p(\mathbf{q}') = \mathbf{v}_{t_n} + \frac{a_p(\mathbf{q}')}{a_{t_n}(\mathbf{q}') + a_p(\mathbf{q}')} \Delta \mathbf{v}.$$

At $\mathbf{q}$, the construction gives $\mathbf{v}_{t_n} + \mathbf{W}_V \mathbf{r} = \bar{\mathbf{v}}_p(\mathbf{q})$. The coefficient of $\Delta \mathbf{v}$ is the logistic function of $\mathbf{q}'^\top \Delta \mathbf{k} / \sqrt{d_k}$. Its derivative is at most $1/4$, so

$$\|\mathbf{v}_{t_n} + \mathbf{W}_V \mathbf{r} - \bar{\mathbf{v}}_p(\mathbf{q}')\|_2 \leq \frac{\|\Delta \mathbf{v}\|_2}{4} \epsilon_k. \quad (10)$$

Next, let w_p be the fraction of attention assigned to the pair in the prompted calculation at $\mathbf{q}'$. The total unnormalized weight of the other positions is identical in both calculations. Holding it fixed, the fraction assigned to the pair or the steered token is a function of its log weight with derivative at most $1/4$. Hence

$$|w_s - w_p| \leq \frac{1}{4} \left| \log \frac{a'_{t_n}(\mathbf{q}')}{a_{t_n}(\mathbf{q}') + a_p(\mathbf{q}')} \right|. \quad (11)$$

The logarithm on the right can be written as

$$\frac{\mathbf{q}'^\top \mathbf{W}_K \mathbf{r}}{\sqrt{d_k}} - \log \left(1 + e^{\mathbf{q}'^\top \Delta \mathbf{k} / \sqrt{d_k}} \right).$$

It is zero at $\mathbf{q}$. The first term changes by at most ϵ_r , and the second by at most ϵ_k , since the derivative of $x \mapsto \log(1 + e^x)$ is between zero and one. Thus

$$|w_s - w_p| \leq \frac{\epsilon_r + \epsilon_k}{4}. \quad (12)$$

To compare the full outputs, first replace the steered value by $\bar{\mathbf{v}}_p(\mathbf{q}')$, keeping all attention weights fixed. This changes the output by at most $w_s \|\mathbf{v}_{t_n} + \mathbf{W}_V \mathbf{r} - \bar{\mathbf{v}}_p(\mathbf{q}')\|_2$. Then change its fraction of attention from w_s to w_p . The relative weights among the other positions are unchanged, so this second change is at most $d_{\text{val}} |w_s - w_p|$: the pair's average and the other positions' average are weighted averages of the original values. The triangle inequality gives

$$\begin{aligned} \|\mathbf{o}_{\text{steer}}(\mathbf{h}') - \mathbf{o}_{\text{prompt}}(\mathbf{h}')\|_2 &\leq w_s \|\mathbf{v}_{t_n} + \mathbf{W}_V \mathbf{r} - \bar{\mathbf{v}}_p(\mathbf{q}')\|_2 + d_{\text{val}} |w_s - w_p| \\ &\leq \frac{w_s \|\Delta \mathbf{v}\|_2 + d_{\text{val}}}{4} \epsilon_k + \frac{d_{\text{val}}}{4} \epsilon_r. \end{aligned} \quad (13)$$

If $\mathcal{I}$ is empty, $w_s = w_p = 1$ and only the first change is needed. Finally, $\mathbf{v}_{t_n} + \mathbf{W}_V \mathbf{r} = \bar{\mathbf{v}}_p(\mathbf{q})$ is itself a weighted average of original values. Both full outputs are therefore weighted averages of that same value set, so their distance is also at most d_{val} . $\square$

Distance to the matching set. Let $\rho = \text{dist}(\mathbf{q}', \mathbf{q} + \mathcal{U}^\perp) = \|P_{\mathcal{U}}(\mathbf{q}' - \mathbf{q})\|_2$, where $P_{\mathcal{U}}$ is the orthogonal projector onto $\mathcal{U}$. Since $\Delta \mathbf{k}$ and $\mathbf{W}_K \mathbf{r}$ lie in $\mathcal{U}$, Cauchy–Schwarz gives

$$\epsilon_k \leq \frac{\|\Delta \mathbf{k}\|_2}{\sqrt{d_k}} \rho, \quad \epsilon_r \leq \frac{\|\mathbf{W}_K \mathbf{r}\|_2}{\sqrt{d_k}} \rho.$$

Using $w_s \leq 1$ and $\|\Delta \mathbf{v}\|_2 \leq d_{\text{val}}$ yields

$$\|\mathbf{o}_{\text{steer}}(\mathbf{h}') - \mathbf{o}_{\text{prompt}}(\mathbf{h}')\|_2 \leq \min \left\{ d_{\text{val}}, \frac{d_{\text{val}}(2\|\Delta \mathbf{k}\|_2 + \|\mathbf{W}_K \mathbf{r}\|_2)}{4\sqrt{d_k}} \rho \right\}.$$

Substituting $\mathbf{q}' = \mathbf{W}_Q \mathbf{h}'$ and $\mathbf{q} = \mathbf{W}_Q \mathbf{h}$ gives the main-text bound, with a coefficient independent of $\mathbf{h}'$.

D.4 Proof of Theorem 2

Theorem 2. Suppose $(\mathbf{W}_Q \mathbf{h})^\top \mathbf{W}_K \mathbf{z} \neq 0$ for some $\mathbf{z} \in \ker(\mathbf{W}_V)$. There exists $\mathbf{r} = \mathbf{r}(\mathbf{h}; \mathbf{H}_S, \mathbf{H}_p, \mathbf{W}_Q, \mathbf{W}_K, \mathbf{W}_V)$ such that steering every position in $\mathcal{S}$ by $\mathbf{r}$ gives

$$\mathbf{o}_{\text{steer}}(\mathbf{h}) = \mathbf{o}_{\text{prompt}}(\mathbf{h}).$$

Write $\mathbf{q} = \mathbf{W}_Q \mathbf{h}$ and $\mathcal{B} = \mathcal{S} \cup \{p_1, \dots, p_m\}$, with the selected positions $\mathcal{S}$ as in the main text. Here s indexes token occurrences, including input, generated, and prompt tokens, and $\mathbf{h}_s$ is the corresponding fixed representation. For either block $\mathcal{E} \in \{\mathcal{S}, \mathcal{B}\}$, define its sum of unnormalized weights and weighted average value by

$$a_s(\mathbf{q}) = e^{\mathbf{q}^\top \mathbf{k}_s / \sqrt{d_k}}, \quad a_{\mathcal{E}}(\mathbf{q}) = \sum_{s \in \mathcal{E}} a_s(\mathbf{q}), \quad \bar{\mathbf{v}}_{\mathcal{E}}(\mathbf{q}) = \frac{\sum_{s \in \mathcal{E}} a_s(\mathbf{q}) \mathbf{v}_s}{a_{\mathcal{E}}(\mathbf{q})}.$$

Both blocks have positive weight sums. Primes denote quantities after steering.

Proof. We keep the common query $\mathbf{q}$ fixed throughout. Adding $\mathbf{r}$ to each representation in $\mathcal{S}$ changes its key by $\mathbf{W}_K \mathbf{r}$ and its value by $\mathbf{W}_V \mathbf{r}$. Thus, for every $s \in \mathcal{S}$,

$$a'_s(\mathbf{q}) = e^{\mathbf{q}^\top \mathbf{W}_K \mathbf{r} / \sqrt{d_k}} a_s(\mathbf{q}), \quad \mathbf{v}'_s = \mathbf{v}_s + \mathbf{W}_V \mathbf{r}.$$

Every weight in $\mathcal{S}$ receives the same multiplier, so the relative weights within the block stay unchanged. Its total weight and weighted average value therefore satisfy

$$a'_{\mathcal{S}}(\mathbf{q}) = e^{\mathbf{q}^\top \mathbf{W}_K \mathbf{r} / \sqrt{d_k}} a_{\mathcal{S}}(\mathbf{q}), \quad \bar{\mathbf{v}}'_{\mathcal{S}}(\mathbf{q}) = \bar{\mathbf{v}}_{\mathcal{S}}(\mathbf{q}) + \mathbf{W}_V \mathbf{r}.$$

We choose $\mathbf{r}$ to make these two quantities equal to $a_{\mathcal{B}}(\mathbf{q})$ and $\bar{\mathbf{v}}_{\mathcal{B}}(\mathbf{q})$, respectively.

Step 1: match the weighted average value. Let $\Delta \bar{\mathbf{h}}$ be the difference between the blocks' attention-weighted mean representations:

$$\Delta \bar{\mathbf{h}} = \frac{\sum_{s \in \mathcal{B}} a_s(\mathbf{q}) \mathbf{h}_s}{a_{\mathcal{B}}(\mathbf{q})} - \frac{\sum_{s \in \mathcal{S}} a_s(\mathbf{q}) \mathbf{h}_s}{a_{\mathcal{S}}(\mathbf{q})}.$$

Since $\mathbf{v}_s = \mathbf{W}_V \mathbf{h}_s$, applying $\mathbf{W}_V$ gives

$$\mathbf{W}_V \Delta \bar{\mathbf{h}} = \bar{\mathbf{v}}_{\mathcal{B}}(\mathbf{q}) - \bar{\mathbf{v}}_{\mathcal{S}}(\mathbf{q}).$$

Thus choosing $\mathbf{r} = \Delta \bar{\mathbf{h}}$ matches the weighted average value. Adding any multiple of $\mathbf{z}$ preserves this match because $\mathbf{W}_V \mathbf{z} = 0$. We use this freedom to adjust the total weight next.

Step 2: match the sum of weights. For $a'_{\mathcal{S}}(\mathbf{q}) = a_{\mathcal{B}}(\mathbf{q})$, the common multiplier must be $a_{\mathcal{B}}(\mathbf{q})/a_{\mathcal{S}}(\mathbf{q})$. Equivalently,

$$\mathbf{q}^\top \mathbf{W}_K \mathbf{r} = \sqrt{d_k} \log \left(\frac{a_{\mathcal{B}}(\mathbf{q})}{a_{\mathcal{S}}(\mathbf{q})} \right).$$

The initial choice $\Delta \bar{\mathbf{h}}$ contributes $\mathbf{q}^\top \mathbf{W}_K \Delta \bar{\mathbf{h}}$ to the left side. Each unit of $\mathbf{z}$ adds $\mathbf{q}^\top \mathbf{W}_K \mathbf{z}$, which is nonzero by assumption. We can therefore supply exactly the remaining amount by setting

$$\mathbf{r} = \Delta \bar{\mathbf{h}} + \frac{\sqrt{d_k} \log \left(\frac{a_{\mathcal{B}}(\mathbf{q})}{a_{\mathcal{S}}(\mathbf{q})} \right) - \mathbf{q}^\top \mathbf{W}_K \Delta \bar{\mathbf{h}}}{\mathbf{q}^\top \mathbf{W}_K \mathbf{z}} \mathbf{z}. \quad (14)$$

This choice preserves Step 1 and satisfies Step 2. Consequently,

$$a'_S(\mathbf{q}) = a_B(\mathbf{q}), \quad a'_S(\mathbf{q})\bar{v}'_S(\mathbf{q}) = a_B(\mathbf{q})\bar{v}_B(\mathbf{q}).$$

The first equality matches the blocks' contributions to the attention denominator; the second matches their contributions to the numerator. All other positions have identical keys and values in the two calculations, so their contributions also agree. Hence

$$\mathbf{o}_{\text{steer}}(\mathbf{h}) = \mathbf{o}_{\text{prompt}}(\mathbf{h}).$$

The same argument applies when there are no other positions.

When $k = 1$, $\mathcal{S} = \{t_n\}$ and the representation difference simplifies to

$$\Delta\bar{\mathbf{h}} = \sum_{j=1}^m \frac{a_{p_j}(\mathbf{q})}{a_B(\mathbf{q})} (\mathbf{h}_{p_j} - \mathbf{h}_{t_n}).$$

With one prompt token, (14) reduces to the single-token construction in (1). $\square$

D.5 Proof of Lemma 3

Lemma 3. Fix the context representations and the displacement $\mathbf{r}$ constructed for Theorem 2 in Appendix D.4. Define

$$\mathcal{U} = \text{col}(\Delta\mathbf{K}_B) + \text{span}\{\mathbf{W}_K\mathbf{z}\}.$$

Then

$$\max\{0, d_k - \text{rank}(\Delta\mathbf{K}_B) - 1\} \leq \dim \mathcal{U}^\perp \leq d_k - \max\{1, \text{rank}(\Delta\mathbf{K}_B)\},$$

and any representation $\mathbf{h}' \in \mathbb{R}^d$ satisfying $\mathbf{W}_Q(\mathbf{h}' - \mathbf{h}) \in \mathcal{U}^\perp$ preserves the match:

$$\mathbf{o}_{\text{steer}}(\mathbf{h}') = \mathbf{o}_{\text{prompt}}(\mathbf{h}').$$

Proof. Write $\mathbf{q} = \mathbf{W}_Q\mathbf{h}$, $\mathbf{q}' = \mathbf{W}_Q\mathbf{h}'$, and $\delta\mathbf{q} = \mathbf{q}' - \mathbf{q}$. At the reference query $\mathbf{q}$, the construction in Appendix D.4 gives

$$\bar{v}_S(\mathbf{q}) + \mathbf{W}_V\mathbf{r} = \bar{v}_B(\mathbf{q}), \quad a'_S(\mathbf{q}) = a_B(\mathbf{q}).$$

We first explain the choice of $\mathcal{U}$, then show that these equalities hold for every $\mathbf{h}'$ satisfying the stated condition.

Why this subspace? The relative weight of a token $s \in \mathcal{B}$ compared with t_n depends on $\mathbf{q}'^\top(\mathbf{k}_s - \mathbf{k}_{t_n})$. Steering multiplies every weight in $\mathcal{S}$ by $e^{\mathbf{q}'^\top \mathbf{W}_K \mathbf{r} / \sqrt{d_k}}$. Thus a sufficient condition for preserving both the relative weights and the steering multiplier is

$$\delta\mathbf{q}^\top(\mathbf{k}_s - \mathbf{k}_{t_n}) = 0 \quad (s \in \mathcal{B}), \quad \delta\mathbf{q}^\top \mathbf{W}_K \mathbf{r} = 0.$$

The key differences are precisely the columns of $\Delta\mathbf{K}_B$. To relate the steering direction to them, recall that $\Delta\bar{\mathbf{h}}$ is the difference of two weighted means. Each mean has weights summing to one, so

$$\mathbf{W}_K \Delta\bar{\mathbf{h}} = \sum_{s \in \mathcal{B}} \frac{a_s(\mathbf{q})}{a_B(\mathbf{q})} (\mathbf{k}_s - \mathbf{k}_{t_n}) - \sum_{s \in \mathcal{S}} \frac{a_s(\mathbf{q})}{a_S(\mathbf{q})} (\mathbf{k}_s - \mathbf{k}_{t_n}).$$

Hence $\mathbf{W}_K \Delta\bar{\mathbf{h}}$ lies in $\text{col}(\Delta\mathbf{K}_B)$. By (14), $\mathbf{r}$ is $\Delta\bar{\mathbf{h}}$ plus a multiple of $\mathbf{z}$. Therefore $\mathbf{W}_K \mathbf{r}$ lies in

$$\mathcal{U} = \text{col}(\Delta\mathbf{K}_B) + \text{span}\{\mathbf{W}_K\mathbf{z}\}.$$

This explains the definition: $\mathcal{U}$ contains all the key differences and the key shift whose inner products with the query must stay fixed. Any $\mathbf{h}'$ with $\mathbf{W}_Q(\mathbf{h}' - \mathbf{h}) \in \mathcal{U}^\perp$ gives a query $\mathbf{q}' = \mathbf{W}_Q\mathbf{h}'$ satisfying both conditions above. We now verify the conclusion for every such $\mathbf{h}'$.

Step 1: preserve the weighted average values. For each $s \in \mathcal{B}$, the first orthogonality condition gives

$$\frac{a_s(\mathbf{q}')}{a_{t_n}(\mathbf{q}')} = e^{\mathbf{q}'^\top(\mathbf{k}_s - \mathbf{k}_{t_n})/\sqrt{d_k}} = e^{\mathbf{q}^\top(\mathbf{k}_s - \mathbf{k}_{t_n})/\sqrt{d_k}} = \frac{a_s(\mathbf{q})}{a_{t_n}(\mathbf{q})}.$$

Thus all tokens in $\mathcal{B}$ retain their relative weights. The same holds within $\mathcal{S} \subset \mathcal{B}$. Since their values are fixed, both blocks retain their weighted average values, and

$$\bar{v}_S(\mathbf{q}') + \mathbf{W}_V\mathbf{r} = \bar{v}_S(\mathbf{q}) + \mathbf{W}_V\mathbf{r} = \bar{v}_B(\mathbf{q}) = \bar{v}_B(\mathbf{q}').$$

Step 2: preserve equality of the total weights. Dividing each block’s weight sum by $a_{t_n}(\mathbf{q}')$ expresses their ratio entirely in terms of the ratios preserved in Step 1:

$$\frac{a_{\mathcal{B}}(\mathbf{q}')}{a_{\mathcal{S}}(\mathbf{q}')} = \frac{\sum_{s \in \mathcal{B}} a_s(\mathbf{q}')/a_{t_n}(\mathbf{q}')}{\sum_{s \in \mathcal{S}} a_s(\mathbf{q}')/a_{t_n}(\mathbf{q}')} = \frac{a_{\mathcal{B}}(\mathbf{q})}{a_{\mathcal{S}}(\mathbf{q})}.$$

Also, $\mathbf{W}_K \mathbf{r} \in \mathcal{U}$ implies $\delta \mathbf{q}^\top \mathbf{W}_K \mathbf{r} = 0$. The steering multiplier therefore satisfies

$$\frac{a'_{\mathcal{S}}(\mathbf{q}')}{a_{\mathcal{S}}(\mathbf{q}')} = e^{\mathbf{q}'^\top \mathbf{W}_K \mathbf{r} / \sqrt{d_k}} = e^{\mathbf{q}^\top \mathbf{W}_K \mathbf{r} / \sqrt{d_k}} = \frac{a_{\mathcal{B}}(\mathbf{q})}{a_{\mathcal{S}}(\mathbf{q})} = \frac{a_{\mathcal{B}}(\mathbf{q}')}{a_{\mathcal{S}}(\mathbf{q}')}.$$

Multiplying by $a_{\mathcal{S}}(\mathbf{q}')$ gives $a'_{\mathcal{S}}(\mathbf{q}') = a_{\mathcal{B}}(\mathbf{q}')$. Together with Step 1, this yields

$$a'_{\mathcal{S}}(\mathbf{q}')(\bar{\mathbf{v}}_{\mathcal{S}}(\mathbf{q}') + \mathbf{W}_V \mathbf{r}) = a_{\mathcal{B}}(\mathbf{q}')\bar{\mathbf{v}}_{\mathcal{B}}(\mathbf{q}').$$

The steered block and the prompted block therefore contribute the same denominator and numerator to attention. All other positions are identical between the two calculations at this query, so

$$\mathbf{o}_{\text{steer}}(\mathbf{h}') = \mathbf{o}_{\text{prompt}}(\mathbf{h}') \quad \text{whenever } \mathbf{W}_Q(\mathbf{h}' - \mathbf{h}) \in \mathcal{U}^\perp.$$

Finally, $\text{col}(\Delta \mathbf{K}_{\mathcal{B}})$ has dimension $\text{rank}(\Delta \mathbf{K}_{\mathcal{B}})$. Adding the direction $\mathbf{W}_K \mathbf{z}$ increases this dimension by at most one. Moreover, $\mathbf{q}^\top \mathbf{W}_K \mathbf{z} \neq 0$ ensures that $\mathbf{W}_K \mathbf{z} \neq 0$, so $\mathcal{U}$ has dimension at least one. Consequently,

$$\max\{1, \text{rank}(\Delta \mathbf{K}_{\mathcal{B}})\} \leq \dim \mathcal{U} \leq \min\{d_k, \text{rank}(\Delta \mathbf{K}_{\mathcal{B}}) + 1\}.$$

Subtracting from d_k gives the stated bounds on $\dim \mathcal{U}^\perp$. $\square$

D.6 Proof of Lemma 4

Lemma 4. Under the same assumptions as Lemma 3, any representation $\mathbf{h}' \in \mathbb{R}^d$ satisfies

$$\|\mathbf{o}_{\text{steer}}(\mathbf{h}') - \mathbf{o}_{\text{prompt}}(\mathbf{h}')\|_2 = O(\text{dist}(\mathbf{W}_Q \mathbf{h}', \mathbf{W}_Q \mathbf{h} + \mathcal{U}^\perp)),$$

where $\mathcal{U} = \text{col}(\Delta \mathbf{K}_{\mathcal{B}}) + \text{span}\{\mathbf{W}_K \mathbf{z}\}$ and dist is Euclidean distance to a set.

Explicit bound. Keep the fixed representations and displacement from Appendix D.5, and write $\mathbf{q} = \mathbf{W}_Q \mathbf{h}$ and $\mathbf{q}' = \mathbf{W}_Q \mathbf{h}'$. Define the two changes in relative attention scores by

$$\epsilon_k = \max_{s \in \mathcal{B}, j \in \mathcal{S}} \frac{|(\mathbf{q}' - \mathbf{q})^\top (\mathbf{k}_s - \mathbf{k}_j)|}{\sqrt{d_k}}, \quad \epsilon_r = \frac{|(\mathbf{q}' - \mathbf{q})^\top \mathbf{W}_K \mathbf{r}|}{\sqrt{d_k}}.$$

The first controls changes in the relative weights within and between the two blocks; the second controls the steering multiplier. Let d_{val} be the largest Euclidean distance between any two values appearing in either calculation, including the shifted values $\mathbf{v}_s + \mathbf{W}_V \mathbf{r}$ for $s \in \mathcal{S}$. Then

$$\|\mathbf{o}_{\text{steer}}(\mathbf{h}') - \mathbf{o}_{\text{prompt}}(\mathbf{h}')\|_2 \leq \min \left\{ d_{\text{val}}, \frac{d_{\text{val}}}{4} (9\epsilon_k + \epsilon_r) \right\}. \quad (15)$$

Proof. The proof follows the same two comparisons as Appendix D.3: changes in weighted average values and changes in attention fractions. In the single-token case, one relative score determines the prompted pair’s average, so the derivative bound for the logistic function controls its change. Here, several relative scores change together, and both block averages can vary. We therefore differentiate along the straight path from $\mathbf{q}$ to $\mathbf{q}'$ and integrate bounds on the derivatives over the entire path. This gives bounds on the finite changes, rather than a first-order approximation.

Step 1: bound the change in weighted average values. Move along the straight segment $\mathbf{q}_\theta = \mathbf{q} + \theta(\mathbf{q}' - \mathbf{q})$, $0 \leq \theta \leq 1$. For either block $\mathcal{E} \in \{\mathcal{S}, \mathcal{B}\}$, differentiating its weighted average gives

$$\frac{d}{d\theta} \bar{\mathbf{v}}_{\mathcal{E}}(\mathbf{q}_\theta) = \sum_{s \in \mathcal{E}} \frac{a_s(\mathbf{q}_\theta)}{a_{\mathcal{E}}(\mathbf{q}_\theta)} (\mathbf{v}_s - \bar{\mathbf{v}}_{\mathcal{E}}(\mathbf{q}_\theta)) \frac{(\mathbf{q}' - \mathbf{q})^\top (\mathbf{k}_s - \mathbf{k}_{t_n})}{\sqrt{d_k}}.$$

Subtracting $\mathbf{k}_{t_n}$ does not change this derivative because the weighted value deviations sum to zero. Each value deviation has norm at most d_{val} , and each score change is at most ϵ_k since $t_n \in \mathcal{S}$. The weights are nonnegative and sum to one, so integrating gives

$$\|\bar{\mathbf{v}}_{\mathcal{E}}(\mathbf{q}') - \bar{\mathbf{v}}_{\mathcal{E}}(\mathbf{q})\|_2 \leq d_{\text{val}}\epsilon_k.$$

At $\mathbf{q}$, the construction matches the two block averages after steering. Applying this bound to both blocks therefore yields

$$\|\bar{\mathbf{v}}_{\mathcal{S}}(\mathbf{q}') + \mathbf{W}_V \mathbf{r} - \bar{\mathbf{v}}_{\mathcal{B}}(\mathbf{q}')\|_2 \leq 2d_{\text{val}}\epsilon_k.$$

Step 2: bound the change in attention fractions. Let w_s and w_p be the fractions of total attention assigned to the steered and prompted blocks at $\mathbf{q}'$. The other positions have identical weights in the two calculations. Holding their total weight fixed, a block's fraction is a function of its log weight with derivative at most $1/4$. Hence

$$|w_s - w_p| \leq \frac{1}{4} \left| \frac{\mathbf{q}'^\top \mathbf{W}_K \mathbf{r}}{\sqrt{d_k}} - \log \left(\frac{a_{\mathcal{B}}(\mathbf{q}')}{a_{\mathcal{S}}(\mathbf{q}')} \right) \right|.$$

The two terms inside the absolute value agree at $\mathbf{q}$. The first changes by at most ϵ_r . For the second,

$$\frac{d}{d\theta} \log \left(\frac{a_{\mathcal{B}}(\mathbf{q}_\theta)}{a_{\mathcal{S}}(\mathbf{q}_\theta)} \right) = \sum_{s \in \mathcal{B}} \sum_{j \in \mathcal{S}} \frac{a_s(\mathbf{q}_\theta) a_j(\mathbf{q}_\theta)}{a_{\mathcal{B}}(\mathbf{q}_\theta) a_{\mathcal{S}}(\mathbf{q}_\theta)} \frac{(\mathbf{q}' - \mathbf{q})^\top (\mathbf{k}_s - \mathbf{k}_j)}{\sqrt{d_k}}.$$

This is a weighted average of score changes, each bounded in absolute value by ϵ_k . Integrating and combining the two changes gives

$$|w_s - w_p| \leq \frac{\epsilon_k + \epsilon_r}{4}.$$

Step 3: combine the two errors. First replace the steered block's weighted average by the prompted block's average, keeping its attention fraction fixed. Then change this fraction from w_s to w_p . The relative weights among all other positions are unchanged. The triangle inequality gives

$$\begin{aligned} \|\mathbf{o}_{\text{steer}}(\mathbf{h}') - \mathbf{o}_{\text{prompt}}(\mathbf{h}')\|_2 &\leq w_s \|\bar{\mathbf{v}}_{\mathcal{S}}(\mathbf{q}') + \mathbf{W}_V \mathbf{r} - \bar{\mathbf{v}}_{\mathcal{B}}(\mathbf{q}')\|_2 + d_{\text{val}} |w_s - w_p| \\ &\leq 2w_s d_{\text{val}} \epsilon_k + \frac{d_{\text{val}}}{4} (\epsilon_k + \epsilon_r) \\ &\leq \frac{d_{\text{val}}}{4} (9\epsilon_k + \epsilon_r), \end{aligned} \tag{16}$$

where the last step uses $w_s \leq 1$. If there are no other positions, $w_s = w_p = 1$ and only the first change is needed. Both complete outputs are weighted averages of values in the set defining d_{val} , so their distance is also at most d_{val} . Unlike the single-token case, matching a block average need not place each shifted value in the convex hull of the original values; this is why the diameter includes shifted values. $\square$

Distance to the matching set. Let $\rho = \text{dist}(\mathbf{q}', \mathbf{q} + \mathcal{U}^\perp)$. Every difference $\mathbf{k}_s - \mathbf{k}_j$, with $s \in \mathcal{B}$ and $j \in \mathcal{S}$, lies in $\text{col}(\Delta \mathbf{K}_{\mathcal{B}})$, and [Appendix D.5](#) shows that $\mathbf{W}_K \mathbf{r} \in \mathcal{U}$. Thus only the component of $\mathbf{q}' - \mathbf{q}$ in $\mathcal{U}$, whose norm is ρ , contributes to the score changes. Cauchy-Schwarz gives

$$\epsilon_k \leq \frac{\max_{s \in \mathcal{B}, j \in \mathcal{S}} \|\mathbf{k}_s - \mathbf{k}_j\|_2}{\sqrt{d_k}} \rho, \quad \epsilon_r \leq \frac{\|\mathbf{W}_K \mathbf{r}\|_2}{\sqrt{d_k}} \rho.$$

Substituting into (15) yields

$$\|\mathbf{o}_{\text{steer}}(\mathbf{h}') - \mathbf{o}_{\text{prompt}}(\mathbf{h}')\|_2 \leq \min \left\{ d_{\text{val}}, \frac{d_{\text{val}} (9 \max_{s \in \mathcal{B}, j \in \mathcal{S}} \|\mathbf{k}_s - \mathbf{k}_j\|_2 + \|\mathbf{W}_K \mathbf{r}\|_2)}{4\sqrt{d_k}} \rho \right\}.$$

The coefficient depends only on the fixed keys, values, projections, and displacement, not on $\mathbf{h}'$. Substituting $\mathbf{q}' = \mathbf{W}_Q \mathbf{h}'$ and $\mathbf{q} = \mathbf{W}_Q \mathbf{h}$ proves the main-text statement.

D.7 Proof of Lemma 5

Here $\mathbf{o}_{\mathcal{S}, \lambda}(\mathbf{h})$ denotes the attention output after adding $\lambda \mathbf{r}$ at every position in $\mathcal{S}$, and $\mathbf{o}_{j, \beta}(\mathbf{h})$ denotes the output after adding $\beta \mathbf{r}$ only at $j \in \mathcal{S}$. Both use the common query $\mathbf{q} = \mathbf{W}_Q \mathbf{h}$ and the same pre-intervention context representations.

Lemma 5. Fix the pre-intervention context representations, projection matrices, and a common input representation $\mathbf{h}$. Compare adding $\lambda \mathbf{r}$ at visible positions $\mathcal{S}$ with adding $\beta \mathbf{r}$ only at $j \in \mathcal{S}$. Using the pre-intervention attention weights $a_i(\mathbf{W}_Q \mathbf{h})$ and setting β as $\beta^* = \lambda \frac{\sum_{i \in \mathcal{S}} a_i(\mathbf{W}_Q \mathbf{h})}{a_j(\mathbf{W}_Q \mathbf{h})}$, then

$$\|\mathbf{o}_{\mathcal{S},\lambda}(\mathbf{h}) - \mathbf{o}_{j,\beta}(\mathbf{h})\|_2 = O\left(\frac{|(\mathbf{W}_Q \mathbf{h})^\top \mathbf{W}_K \mathbf{r}|}{\sqrt{d_k}}\right).$$

Explicit bound. Set $\beta = \beta^*$ as in the lemma, and let $\mathcal{C}$ be the common set of visible positions. For each position $i \in \mathcal{C}$, recall its original key $\mathbf{k}_i$, value $\mathbf{v}_i$, and unnormalized weight $a_i(\mathbf{q}) = e^{\mathbf{q}^\top \mathbf{k}_i / \sqrt{d_k}}$. Define

$$c = \frac{\mathbf{q}^\top \mathbf{W}_K \mathbf{r}}{\sqrt{d_k}}, \quad d_{\text{val}} = \max_{i,l \in \mathcal{C}} \|\mathbf{v}_i - \mathbf{v}_l\|_2.$$

Here c is the attention-score change from adding $\mathbf{r}$ once, and d_{val} measures the spread of the original values. Then

$$\|\mathbf{o}_{\mathcal{S},\lambda}(\mathbf{h}) - \mathbf{o}_{j,\beta}(\mathbf{h})\|_2 \leq \frac{|c|}{4} [d_{\text{val}}(|\lambda| + |\beta|) + (\lambda^2 + \beta^2) \|\mathbf{W}_V \mathbf{r}\|_2]. \quad (17)$$

Proof. The choice of β makes the two interventions agree if attention weights are held at their original values. We first show this equality, then bound the error caused by changes in those weights.

Step 1: match the value additions at the original weights. For either selected set $\mathcal{E} = \mathcal{S}$ or $\mathcal{E} = \{j\}$, let

$$w_{\mathcal{E}} = \frac{\sum_{i \in \mathcal{E}} a_i(\mathbf{q})}{\sum_{l \in \mathcal{C}} a_l(\mathbf{q})}$$

be its original fraction of attention, and let $\mathbf{o}_0$ be the attention output before either intervention. Holding these weights fixed, the two interventions would give

$$\mathbf{o}_0 + \lambda w_{\mathcal{S}} \mathbf{W}_V \mathbf{r}, \quad \mathbf{o}_0 + \beta w_{\{j\}} \mathbf{W}_V \mathbf{r},$$

respectively: each shifted value contributes its original attention weight times the added value. The chosen strength satisfies

$$\beta w_{\{j\}} = \lambda \frac{\sum_{i \in \mathcal{S}} a_i(\mathbf{q})}{a_j(\mathbf{q})} \frac{a_j(\mathbf{q})}{\sum_{l \in \mathcal{C}} a_l(\mathbf{q})} = \lambda w_{\mathcal{S}}.$$

Thus both interventions have the same output when evaluated with the original weights.

Step 2: bound the effect of changing attention weights. We use one calculation for both interventions: write $(\mathcal{E}, \gamma) = (\mathcal{S}, \lambda)$ for steering the selected block, or $(\mathcal{E}, \gamma) = (\{j\}, \beta)$ for steering only j . Thus γ denotes the strength in the case under consideration, and $\mathbf{o}_{\mathcal{E},\gamma}(\mathbf{h})$ denotes its actual attention output. Adding $\gamma \mathbf{r}$ multiplies every weight in $\mathcal{E}$ by $e^{\gamma c}$. The relative weights inside the block and among the other positions remain unchanged; only the block's total fraction of attention changes. Denote its fraction after steering by $w'_{\mathcal{E}}$; then

$$w'_{\mathcal{E}} = \frac{w_{\mathcal{E}} e^{\gamma c}}{1 - w_{\mathcal{E}} + w_{\mathcal{E}} e^{\gamma c}}.$$

For $0 < w_{\mathcal{E}} < 1$, this is a logistic function of γc . As in the preceding approximation proofs, its derivative is at most $1/4$, giving

$$|w'_{\mathcal{E}} - w_{\mathcal{E}}| \leq \frac{|\gamma c|}{4}.$$

When $\mathcal{E} \neq \mathcal{C}$, compare the actual output with $\mathbf{o}_0 + \gamma w_{\mathcal{E}} \mathbf{W}_V \mathbf{r}$ by keeping the shifted values fixed and changing the block's attention fraction from $w_{\mathcal{E}}$ to $w'_{\mathcal{E}}$. The original weighted average values inside and outside the block differ by at most d_{val} . Steering adds $\gamma \mathbf{W}_V \mathbf{r}$ to the first average, so the difference after steering is at most $d_{\text{val}} + |\gamma| \|\mathbf{W}_V \mathbf{r}\|_2$. Consequently,

$$\begin{aligned} \|\mathbf{o}_{\mathcal{E},\gamma}(\mathbf{h}) - (\mathbf{o}_0 + \gamma w_{\mathcal{E}} \mathbf{W}_V \mathbf{r})\|_2 &\leq |w'_{\mathcal{E}} - w_{\mathcal{E}}| (d_{\text{val}} + |\gamma| \|\mathbf{W}_V \mathbf{r}\|_2) \\ &\leq \frac{|\gamma c|}{4} (d_{\text{val}} + |\gamma| \|\mathbf{W}_V \mathbf{r}\|_2). \end{aligned}$$

If $\mathcal{E} = \mathcal{C}$, its fraction remains one and the output is exactly $\mathbf{o}_0 + \gamma \mathbf{W}_V \mathbf{r}$, so the same bound holds.

Step 3: compare the two interventions. Step 1 gives $\lambda w_{\mathcal{S}} = \beta w_{\{j\}}$, so both outputs computed with the original weights equal $\mathbf{o}_0 + \lambda w_{\mathcal{S}} \mathbf{W}_V \mathbf{r}$. Insert this common reference output between the two actual outputs and apply the triangle inequality:

$$\begin{aligned} \|\mathbf{o}_{\mathcal{S},\lambda}(\mathbf{h}) - \mathbf{o}_{j,\beta}(\mathbf{h})\|_2 &\leq \|\mathbf{o}_{\mathcal{S},\lambda}(\mathbf{h}) - (\mathbf{o}_0 + \lambda w_{\mathcal{S}} \mathbf{W}_V \mathbf{r})\|_2 + \|\mathbf{o}_{j,\beta}(\mathbf{h}) - (\mathbf{o}_0 + \lambda w_{\mathcal{S}} \mathbf{W}_V \mathbf{r})\|_2 \\ &= \|\mathbf{o}_{\mathcal{S},\lambda}(\mathbf{h}) - (\mathbf{o}_0 + \lambda w_{\mathcal{S}} \mathbf{W}_V \mathbf{r})\|_2 + \|\mathbf{o}_{j,\beta}(\mathbf{h}) - (\mathbf{o}_0 + \beta w_{\{j\}} \mathbf{W}_V \mathbf{r})\|_2. \end{aligned}$$

The two terms now have exactly the form bounded in Step 2. Apply that bound with $(\mathcal{E}, \gamma) = (\mathcal{S}, \lambda)$ to the first and $(\mathcal{E}, \gamma) = (\{j\}, \beta)$ to the second:

$$\begin{aligned} \|\mathbf{o}_{\mathcal{S},\lambda}(\mathbf{h}) - \mathbf{o}_{j,\beta}(\mathbf{h})\|_2 &\leq \frac{|\lambda c|}{4} (d_{\text{val}} + |\lambda| \|\mathbf{W}_V \mathbf{r}\|_2) + \frac{|\beta c|}{4} (d_{\text{val}} + |\beta| \|\mathbf{W}_V \mathbf{r}\|_2) \\ &= \frac{|c|}{4} [|\lambda| d_{\text{val}} + \lambda^2 \|\mathbf{W}_V \mathbf{r}\|_2 + |\beta| d_{\text{val}} + \beta^2 \|\mathbf{W}_V \mathbf{r}\|_2] \\ &= \frac{|c|}{4} [d_{\text{val}} (|\lambda| + |\beta|) + (\lambda^2 + \beta^2) \|\mathbf{W}_V \mathbf{r}\|_2], \end{aligned}$$

which is (17). Since $c = (\mathbf{W}_Q \mathbf{h})^\top \mathbf{W}_K \mathbf{r} / \sqrt{d_k}$, this gives the main-text conclusion. The coefficient multiplying $|c|$ depends on the fixed values, displacement, and the two strengths, including the attention-dependent β ; it need not be uniform if β is recalibrated at different queries. $\square$

When $c = 0$, neither intervention changes the attention weights, so the outputs agree exactly. If the $k = |\mathcal{S}|$ selected positions have equal original weights, the strength reduces to $\beta = k\lambda$. In general, the required strength depends on their attention weights, not just their number: with two positions of attention fractions $1/4$ and $3/4$, matching the value additions at those weights using only the first position requires $\beta = 4\lambda$.

E Multi-Head Attention-Output Matching

E.1 Multi-head output matching

We extend the single-token comparison to H attention heads at one fixed layer. The same displacement $\mathbf{r}$ is added to $\mathbf{h}_{t_n}$ in every head, while prompting adds token p . We seek equality of the combined outputs after the head-specific output projections, using the same residual input $\mathbf{h}$.

Notation. Fix the projections and pre-intervention representations. For head s , let $\mathbf{q}^{(s)} = \mathbf{W}_Q^{(s)} \mathbf{h}$ and $\mathbf{W}_O^{(s)} \in \mathbb{R}^{d \times d_v}$ be its query and output projection. Each token occurrence j has original weight and value

$$a_j^{(s)} = \exp\left(\frac{(\mathbf{q}^{(s)})^\top \mathbf{W}_K^{(s)} \mathbf{h}_j}{\sqrt{d_k}}\right), \quad \mathbf{v}_j^{(s)} = \mathbf{W}_V^{(s)} \mathbf{h}_j.$$

The shared positions $\mathcal{I} = \{t_1, \dots, t_{n-1}\}$ have identical keys and values in both calculations. For each head s , let η_s be a proposed change in the attention score of t_n , and collect these changes in

$$\boldsymbol{\eta} = (\eta_1, \dots, \eta_H)^\top \in \mathbb{R}^H.$$

These are candidate changes until a displacement realizing them is found. Define the score map $\mathbf{S}_h \in \mathbb{R}^{H \times d}$ and the resulting attention fraction $w_s(\eta_s)$ by

$$\mathbf{S}_h[s, :] = \frac{(\mathbf{q}^{(s)})^\top \mathbf{W}_K^{(s)}}{\sqrt{d_k}}, \quad w_s(\eta_s) = \frac{a_{t_n}^{(s)} e^{\eta_s}}{\sum_{j \in \mathcal{I}} a_j^{(s)} + a_{t_n}^{(s)} e^{\eta_s}}.$$

Two quantities collect the projected contributions of the added values and the original values, respectively:

$$\begin{aligned} \mathbf{M}_h(\boldsymbol{\eta}) &= \sum_{s=1}^H w_s(\eta_s) \mathbf{W}_O^{(s)} \mathbf{W}_V^{(s)}, \\ \mathbf{c}_h(\boldsymbol{\eta}) &= \sum_{s=1}^H \mathbf{W}_O^{(s)} \frac{\sum_{j \in \mathcal{I}} a_j^{(s)} \mathbf{v}_j^{(s)} + a_{t_n}^{(s)} e^{\eta_s} \mathbf{v}_{t_n}^{(s)}}{\sum_{j \in \mathcal{I}} a_j^{(s)} + a_{t_n}^{(s)} e^{\eta_s}}. \end{aligned} \tag{18}$$

Here $M_h(\eta)$ maps a displacement to its added output contribution, whereas $c_h(\eta)$ uses the original values with the changed weights. The proof derives this decomposition. Write $\mathbf{o}_{\text{steer}}^{(s)}(\mathbf{h})$ and $\mathbf{o}_{\text{prompt}}^{(s)}(\mathbf{h})$ for the two head outputs, and $\tilde{\mathbf{h}}_{\text{steer}}(\mathbf{h})$ and $\tilde{\mathbf{h}}_{\text{prompt}}(\mathbf{h})$ for their projected sums plus $\mathbf{h}$.

Lemma 6 (Projected multi-head matching). *Fix the projection matrices, pre-intervention context representations, and common representation $\mathbf{h}$. A displacement $\mathbf{r}$ matches the complete attention residual update if and only if there is $\eta \in \mathbb{R}^H$ such that*

$$\begin{aligned} S_h \mathbf{r} &= \eta, \\ M_h(\eta) \mathbf{r} &= \sum_{s=1}^H \mathbf{W}_O^{(s)} \mathbf{o}_{\text{prompt}}^{(s)}(\mathbf{h}) - c_h(\eta). \end{aligned} \quad (19)$$

The first equation requires the same $\mathbf{r}$ to realize every proposed score change. The second requires its added values to supply the remaining difference from the prompted output. Together they are equivalent to

$$\tilde{\mathbf{h}}_{\text{steer}}(\mathbf{h}) = \tilde{\mathbf{h}}_{\text{prompt}}(\mathbf{h}).$$

Proof. We first determine the changes produced by a given $\mathbf{r}$, then derive the combined output. This will show why both equations are needed and why they suffice.

Step 1: determine the score changes produced by the same displacement. In head s , adding $\mathbf{r}$ changes the key of t_n by $\mathbf{W}_K^{(s)} \mathbf{r}$. Its attention score therefore changes by

$$\frac{(\mathbf{q}^{(s)})^\top \mathbf{W}_K^{(s)} (\mathbf{h}_{t_n} + \mathbf{r})}{\sqrt{d_k}} - \frac{(\mathbf{q}^{(s)})^\top \mathbf{W}_K^{(s)} \mathbf{h}_{t_n}}{\sqrt{d_k}} = \frac{(\mathbf{q}^{(s)})^\top \mathbf{W}_K^{(s)} \mathbf{r}}{\sqrt{d_k}} = S_h[s, :] \mathbf{r}.$$

Consequently, proposed score changes η are realized by $\mathbf{r}$ exactly when $S_h \mathbf{r} = \eta$. This is the first equation in (19). In the next two steps, take η to be these actual changes.

Step 2: separate each head's original and added values. The score change η_s multiplies the weight of t_n by e^{η_s} . Its value changes from $\mathbf{v}_{t_n}^{(s)}$ to $\mathbf{v}_{t_n}^{(s)} + \mathbf{W}_V^{(s)} \mathbf{r}$, while all other keys and values remain fixed. Thus

$$\begin{aligned} \mathbf{o}_{\text{steer}}^{(s)}(\mathbf{h}) &= \frac{\sum_{j \in \mathcal{I}} a_j^{(s)} \mathbf{v}_j^{(s)} + a_{t_n}^{(s)} e^{\eta_s} (\mathbf{v}_{t_n}^{(s)} + \mathbf{W}_V^{(s)} \mathbf{r})}{\sum_{j \in \mathcal{I}} a_j^{(s)} + a_{t_n}^{(s)} e^{\eta_s}} \\ &= \frac{\sum_{j \in \mathcal{I}} a_j^{(s)} \mathbf{v}_j^{(s)} + a_{t_n}^{(s)} e^{\eta_s} \mathbf{v}_{t_n}^{(s)}}{\sum_{j \in \mathcal{I}} a_j^{(s)} + a_{t_n}^{(s)} e^{\eta_s}} + w_s(\eta_s) \mathbf{W}_V^{(s)} \mathbf{r}. \end{aligned}$$

The first term uses the original values with the new attention weights. The second is the added value multiplied by the fraction of attention assigned to t_n . Both terms are well-defined even if $\mathcal{I}$ is empty.

Step 3: combine the heads and compare their complete outputs. We still use the actual score changes $\eta = S_h \mathbf{r}$ from Step 1. Multiply each head output in Step 2 by $\mathbf{W}_O^{(s)}$ and sum. Keeping the original-value terms and added-value terms separate gives

$$\begin{aligned} \sum_{s=1}^H \mathbf{W}_O^{(s)} \mathbf{o}_{\text{steer}}^{(s)}(\mathbf{h}) &= \sum_{s=1}^H \mathbf{W}_O^{(s)} \frac{\sum_{j \in \mathcal{I}} a_j^{(s)} \mathbf{v}_j^{(s)} + a_{t_n}^{(s)} e^{\eta_s} \mathbf{v}_{t_n}^{(s)}}{\sum_{j \in \mathcal{I}} a_j^{(s)} + a_{t_n}^{(s)} e^{\eta_s}} \\ &\quad + \sum_{s=1}^H w_s(\eta_s) \mathbf{W}_O^{(s)} \mathbf{W}_V^{(s)} \mathbf{r} \\ &= c_h(\eta) + \left(\sum_{s=1}^H w_s(\eta_s) \mathbf{W}_O^{(s)} \mathbf{W}_V^{(s)} \right) \mathbf{r} \\ &= c_h(\eta) + M_h(\eta) \mathbf{r}. \end{aligned}$$

The original-value sum is exactly the definition of $c_h(\eta)$. In the added-value sum, the same $\mathbf{r}$ appears in every term, so we can factor it out; its coefficient is $M_h(\eta)$. Adding the common residual input gives

$$\begin{aligned} \tilde{\mathbf{h}}_{\text{steer}}(\mathbf{h}) &= \mathbf{h} + c_h(\eta) + M_h(\eta) \mathbf{r}, \\ \tilde{\mathbf{h}}_{\text{prompt}}(\mathbf{h}) &= \mathbf{h} + \sum_{s=1}^H \mathbf{W}_O^{(s)} \mathbf{o}_{\text{prompt}}^{(s)}(\mathbf{h}). \end{aligned}$$

Subtract the second update from the first. Their common $\mathbf{h}$ cancels, leaving

$$\tilde{\mathbf{h}}_{\text{steer}}(\mathbf{h}) - \tilde{\mathbf{h}}_{\text{prompt}}(\mathbf{h}) = \mathbf{c}_h(\boldsymbol{\eta}) + \mathbf{M}_h(\boldsymbol{\eta})\mathbf{r} - \sum_{s=1}^H \mathbf{W}_O^{(s)} \mathbf{o}_{\text{prompt}}^{(s)}(\mathbf{h}).$$

Thus the updates agree exactly when the right side is zero. Moving the original-value contribution to the other side gives

$$\mathbf{M}_h(\boldsymbol{\eta})\mathbf{r} = \sum_{s=1}^H \mathbf{W}_O^{(s)} \mathbf{o}_{\text{prompt}}^{(s)}(\mathbf{h}) - \mathbf{c}_h(\boldsymbol{\eta}),$$

which is the second equation in (19).

Step 4: prove necessity and sufficiency. For necessity, suppose a displacement $\mathbf{r}$ gives equal complete updates. Choose $\boldsymbol{\eta} = \mathbf{S}_h \mathbf{r}$, the vector of actual score changes derived in Step 1. This choice satisfies the first equation in (19). Steps 2 and 3 then apply to the actual head outputs. Their difference is zero by assumption, so the last equation of Step 3 holds as well. Hence both equations are necessary.

For sufficiency, suppose a displacement $\mathbf{r}$ and a candidate vector $\boldsymbol{\eta} \in \mathbb{R}^H$ satisfy both equations in (19). The first equation, $\mathbf{S}_h \mathbf{r} = \boldsymbol{\eta}$, ensures that each candidate η_s is the score change actually produced by $\mathbf{r}$. Therefore the weights in Step 2 and the combined output in Step 3 are the actual ones. Substituting the second equation into that combined output gives

$$\begin{aligned} \tilde{\mathbf{h}}_{\text{steer}}(\mathbf{h}) &= \mathbf{h} + \mathbf{c}_h(\boldsymbol{\eta}) + \mathbf{M}_h(\boldsymbol{\eta})\mathbf{r} \\ &= \mathbf{h} + \mathbf{c}_h(\boldsymbol{\eta}) + \sum_{s=1}^H \mathbf{W}_O^{(s)} \mathbf{o}_{\text{prompt}}^{(s)}(\mathbf{h}) - \mathbf{c}_h(\boldsymbol{\eta}) \\ &= \mathbf{h} + \sum_{s=1}^H \mathbf{W}_O^{(s)} \mathbf{o}_{\text{prompt}}^{(s)}(\mathbf{h}) \\ &= \tilde{\mathbf{h}}_{\text{prompt}}(\mathbf{h}). \end{aligned}$$

This proves sufficiency and completes the equivalence. $\square$

The examples in [Appendix E.2](#) explain when these conditions admit a common displacement and how output projection affects matching.

E.2 Examples illustrating multi-head matching

The following examples clarify three aspects of Lemma 6. First, the single-head construction’s two conditions, matching total weight and weighted average value, are sufficient but need not be necessary for equal outputs. Second, even when each head can be matched separately, their required displacements may be incompatible. Third, output projection can allow a match by canceling differences between heads. We use two heads throughout and write $\mathbf{W}_O = [\mathbf{W}_O^{(1)} \cdots \mathbf{W}_O^{(H)}]$ for the combined output projection. The first two examples take the output projection to be the identity, so each head must match; the third changes only this projection to demonstrate cancellation.

Example 1: the construction conditions are sufficient, not necessary. We first show that matching an attention output does not require separately matching the total weight and weighted average value used in our construction. The output is a normalized weighted sum, so different weights and values can produce the same result. We choose two heads for which a common displacement matches both outputs, then show why that displacement cannot satisfy both construction conditions. Take $H = d = 2$, $d_k = d_v = 1$, $\mathbf{h} = (1, 0)^\top$, and one shared context representation $(0, 0)^\top$. Set $\mathbf{W}_Q^{(1)} = \mathbf{W}_Q^{(2)} = (1, 0)$ and

$$\mathbf{W}_V^{(1)} = (1, 0), \quad \mathbf{W}_V^{(2)} = (0, 1), \quad \mathbf{W}_K^{(1)} = (0, 1), \quad \mathbf{W}_K^{(2)} = (1, 0). \quad (20)$$

Let $\mathbf{W}_O$ be the identity, $\mathbf{h}_{t_n} = (0, 0)^\top$, and $\mathbf{h}_p = (1, 1)^\top$. Each prompted head output is $e/(2 + e)$. Steering with $\mathbf{r} = (t, t)^\top$ gives $te^t/(1 + e^t)$ in both heads. This function is strictly increasing for $t > 0$, starts at zero, and is unbounded, so there is a unique positive t matching the prompted output. By contrast, making the steered value equal to the prompted pair’s weighted average in each head forces

both coordinates of $\mathbf{r}$ to equal $e/(1+e)$. Making the steered weight equal to the sum of the pair's weights instead forces both coordinates to equal $\log(1+e)$. Those stronger constraints are inconsistent. The projected-output criterion therefore admits an exact match even though the two sufficient conditions used in the single-head construction cannot hold in both heads at once.

Example 2: separate head-wise solutions need not give a common solution. The previous example admits one displacement that works for both heads, but Lemma 6 does not guarantee such a solution. For fixed proposed score changes $\boldsymbol{\eta}$, the matrices in (19) are known, and its two equations form a linear system for the same unknown $\mathbf{r}$. A solution must both realize the proposed score changes and match the projected output. An exact match exists only if this joint system is solvable for some $\boldsymbol{\eta}$; failure at one choice of $\boldsymbol{\eta}$ does not rule out all others.

The single-head theorem does not resolve this question either: the displacement it provides can differ from one head to another. We now change the token representations so that each head remains individually matchable, but no displacement works for both. We prove this by deriving a contradiction from the two required output equalities.

Keep the same projections, representation $\mathbf{h}$, and shared zero representation, but set $\mathbf{h}_{t_n} = \mathbf{h}_p = (-1, -1)^\top$. Both prompted head outputs now equal $-c$, where $c = 2/(e+2)$. For an arbitrary steered representation $\mathbf{h}_{t_n} + \mathbf{r} = (x, y)^\top$, the two outputs are

$$\mathbf{o}_{\text{steer}}^{(1)} = \frac{xe^y}{1+e^y}, \quad \mathbf{o}_{\text{steer}}^{(2)} = \frac{ye^x}{1+e^x}. \quad (21)$$

Because $\mathbf{W}_O$ is the identity, a match requires both outputs to equal $-c$. Hence $x, y < 0$. Put $u = -x$ and $v = -y$. The matching equations imply

$$\frac{u}{1+e^v} = \frac{v}{1+e^u} = c, \quad u(1+e^u) = v(1+e^v). \quad (22)$$

The function $u \mapsto u(1+e^u)$ is strictly increasing for $u > 0$, so $u = v$. But then

$$c = \frac{u}{1+e^u} < ue^{-u} \leq \frac{1}{e} < \frac{2}{e+2} = c, \quad (23)$$

which is impossible. Each head nevertheless satisfies the assumption of Theorem 1: $\mathbf{z} = (0, 1)^\top$ works for head 1 and $\mathbf{z} = (1, 0)^\top$ for head 2. Thus even individual single-head existence does not ensure a common displacement matching the final projected output.

Example 3: projection can cancel differences between heads. The impossibility in Example 2 relies on the identity output projection, which requires each head to match. More generally, the final update depends on the sum of projected head outputs. Differences between heads may therefore cancel after projection. Stack these differences in

$$\mathbf{e} = \begin{bmatrix} \mathbf{o}_{\text{steer}}^{(1)}(\mathbf{h}) - \mathbf{o}_{\text{prompt}}^{(1)}(\mathbf{h}) \\ \vdots \\ \mathbf{o}_{\text{steer}}^{(H)}(\mathbf{h}) - \mathbf{o}_{\text{prompt}}^{(H)}(\mathbf{h}) \end{bmatrix}.$$

The difference between the complete residual updates is $\mathbf{W}_O \mathbf{e}$. Thus they agree precisely when

$$\mathbf{W}_O \mathbf{e} = \mathbf{0}.$$

If $\mathbf{W}_O$ has full column rank, this forces $\mathbf{e} = \mathbf{0}$, so every head must match. Otherwise, a nonzero $\mathbf{e}$ in its nullspace is also allowed: projection removes the combined effect of those head differences. We illustrate this by keeping the head computations from Example 2 and changing only the output projection.

In the preceding example, replace the output projection by $\mathbf{W}_O = \begin{bmatrix} 1 & 1 \\ 0 & 0 \end{bmatrix}$. Choose the steered representation $(x, y) = (-4c, 0)$, hence $\mathbf{r} = (1-4c, 1)^\top$. Its two head outputs are $(-2c, 0)$, while the prompted head outputs are $(-c, -c)$. Their difference $(-c, c)^\top$ is nonzero but lies in $\ker \mathbf{W}_O$, so the projected outputs and residual updates agree exactly. This illustrates the cancellation allowed by Lemma 6.